

# Internet Protocols EBU5403: The Transport layer - Part I

Michael Chai (michael.chai@qmul.ac.uk)

Richard Clegg (r.clegg@qmul.ac.uk)

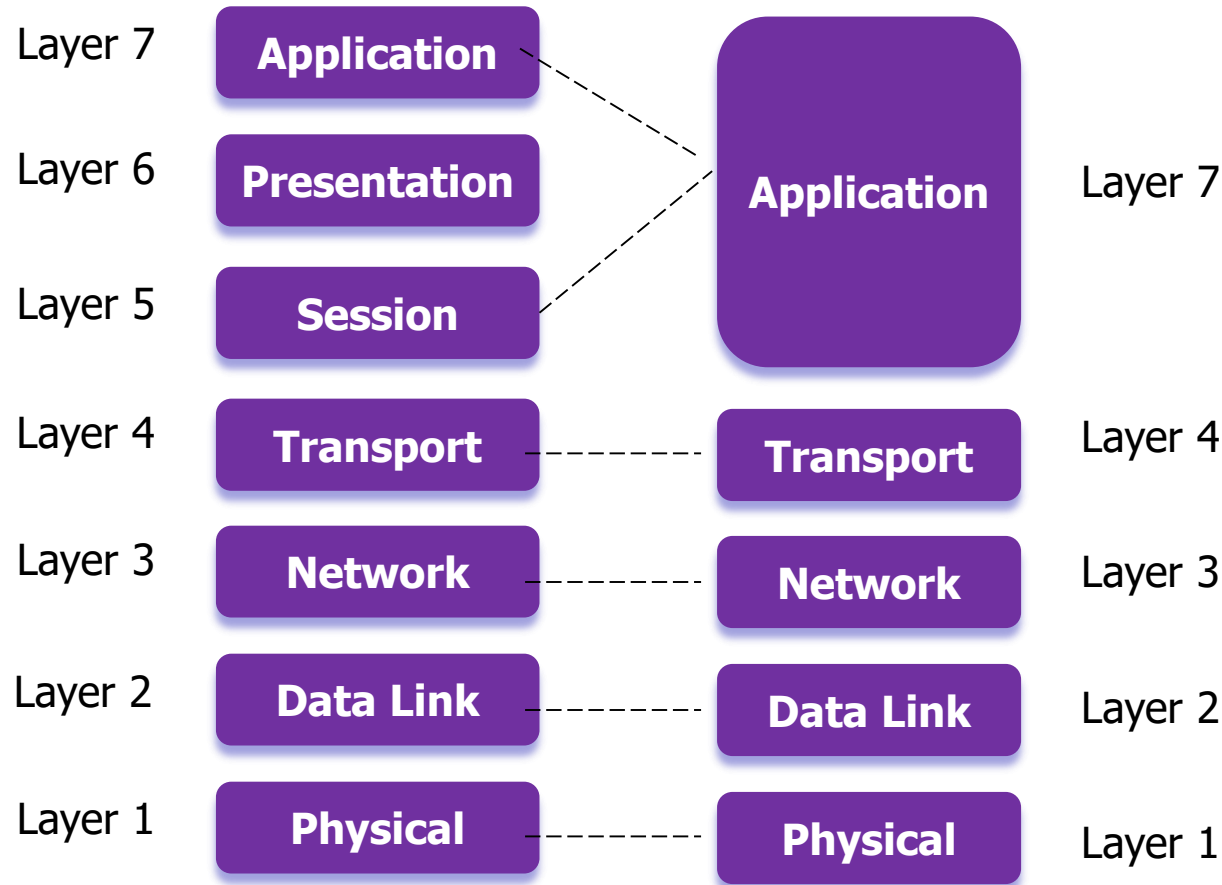
Adnan Kiani (adnankhal@gmail.com)

	Week 1	Week 2	Week 3	Week 4
Telecom	Adnan Kiani		Michael Chai	
E-Commerce	Richard Clegg			

# Structure of course

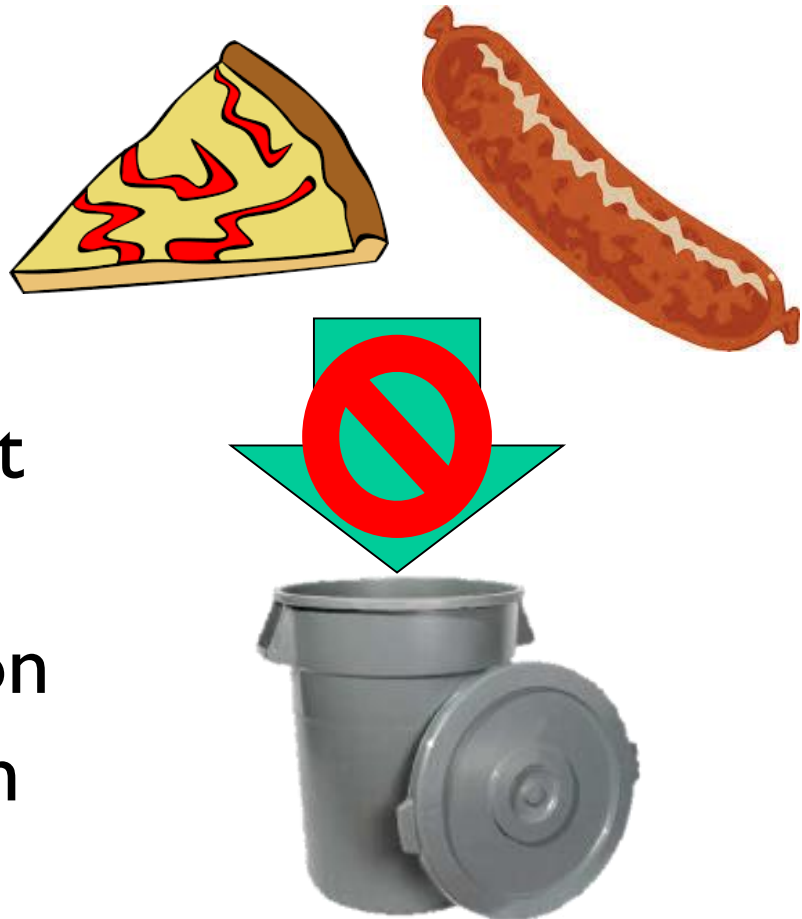
- Week 1 (8<sup>th</sup>-12<sup>th</sup> October)
  - Introduction to IP Networks
  - The Transport layer (part I)
- Week 2 (29<sup>th</sup> October - 2<sup>nd</sup> November)
  - The Transport layer (part II)
  - The Network layer (part I)
  - Class test
- Week 3 (19<sup>th</sup>-23<sup>rd</sup> November)
  - The Network layer (part II)
  - The Data link layer (part I)
  - Router lab tutorial (assessed lab work after this week)
- Week 4 (10<sup>th</sup>-14<sup>th</sup> December)
  - The Data link layer (part II)
  - Network management and security
  - Class test

# ISO/OSI (left) vs TCP/IP (right)



# How to remember the layers

- Please Do Not Throw Pizza + Sausage Away
- **P**lease – **P**hysical
- **D**o – **D**atalink
- **N**ot – **N**etwork
- **T**hrow – **T**ransport
- **S**ausage – **S**ession
- **P**izza – **P**resentation
- **A**way – **A**pplication

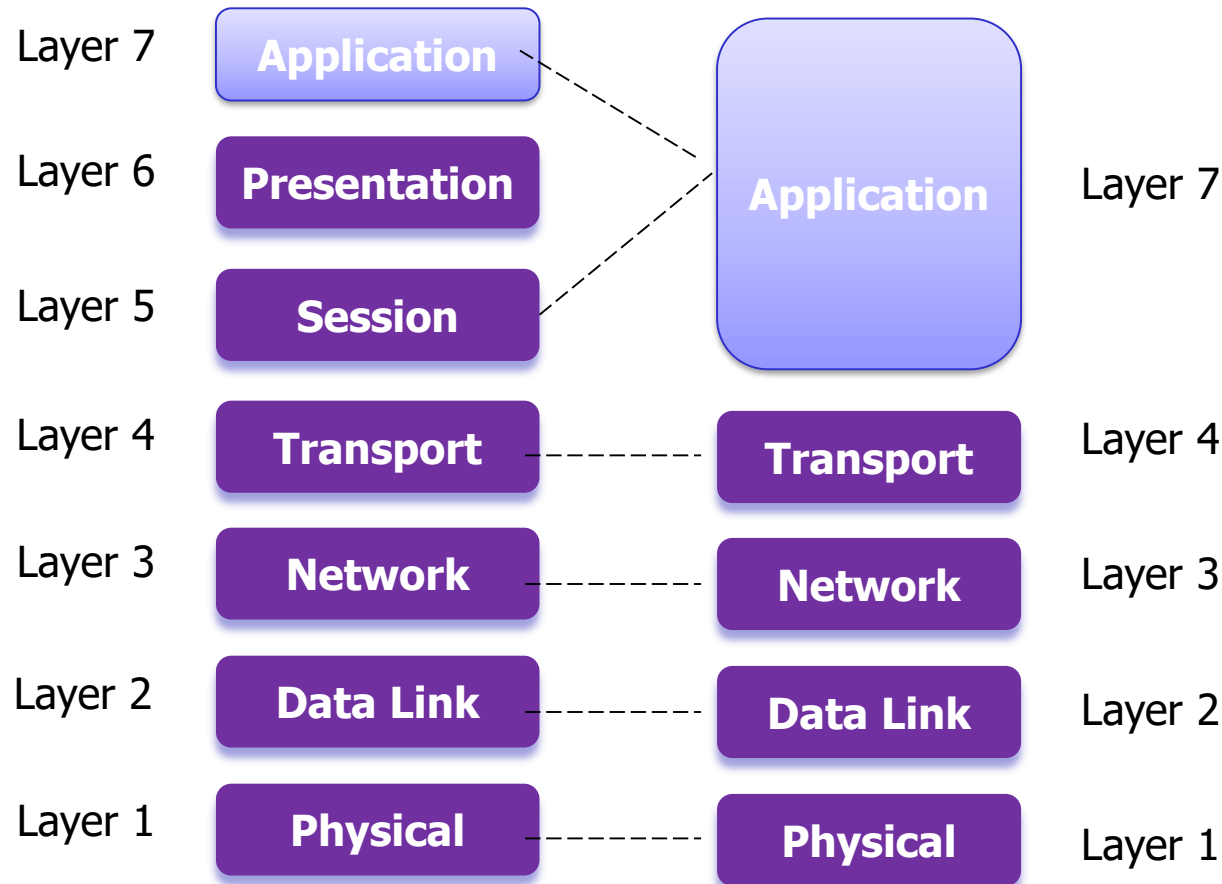


# Week 1: Transport Layer (part 1)

## our goals:

- understand principles behind transport layer services:
  - multiplexing, demultiplexing
  - reliable data transfer
  - flow control
  - congestion control
- learn about Internet transport layer protocols:
  - UDP: connectionless transport
  - TCP: connection-oriented reliable transport
  - TCP congestion control

# Application layer (very briefly)

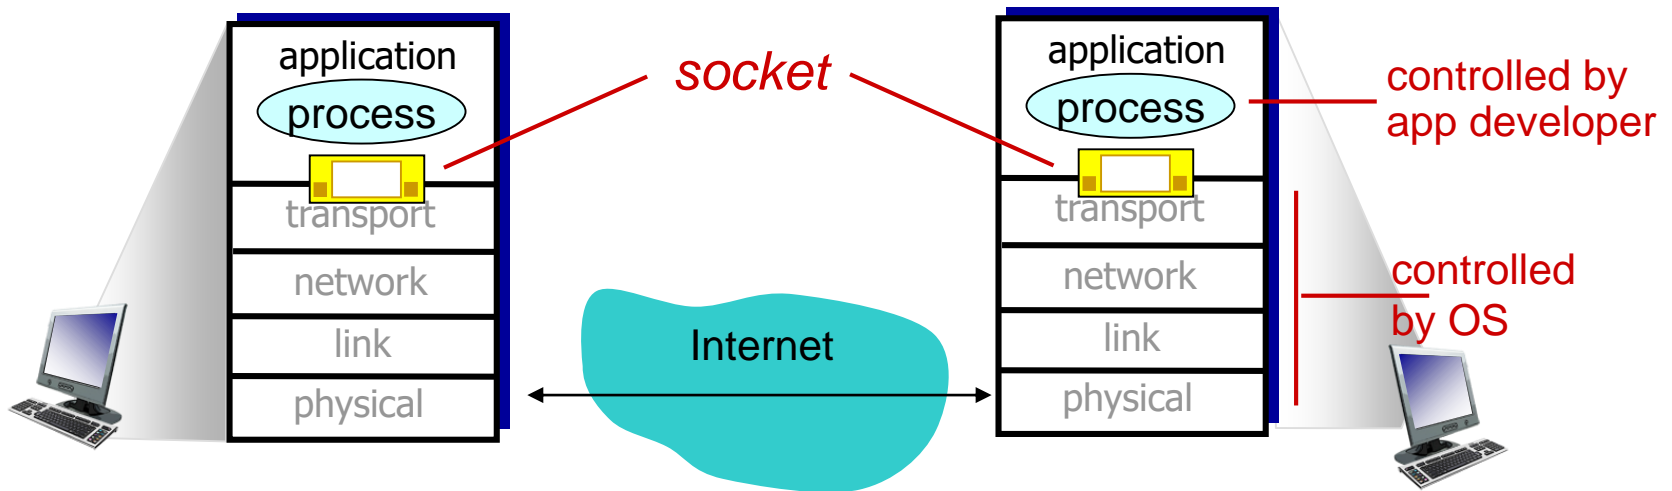


# Application Layer (very briefly)

- This is covered in many other modules you will take.
- This is the layer you (mostly) work with as a programmer.
- At the application layer the “network” is abstracted away and you access a stream of data that arrives at a “socket”.
- It is up to you to define the format of data your application writes and receives.
- Different applications have different formats:
  - HTTP (Hypertext transfer protocol) world-wide web
  - FTP (File transfer protocol) moving data
  - SMTP (Send Mail transfer protocol) sending email
  - IMAP (Internet message access protocol) receiving email

# Sockets

- process sends/receives messages to/from its **socket**
- socket analogous to door
  - sending process shoves message out door
  - sending process relies on transport infrastructure on other side of door to deliver message to socket at receiving process





# Addressing processes

- to receive messages, process must have *identifier*
- host device has unique 32-bit IP address
- Q: does IP address of host on which process runs suffice for identifying the process?
  - A: no, *many* processes can be running on same host
- *identifier* includes both **IP address** and **port numbers** associated with process on host.
- example port numbers:
  - HTTP server: 80
  - mail server: 25
- to send HTTP message to gaia.cs.umass.edu web server:
  - **IP address**: 128.119.245.12
  - **port number**: 80
- more shortly...

# 4 different addresses in TCP/IP

- Physical address Layer 2
  - Also known as the link address, is the address of a node as defined by its LAN or WAN
- Logical address Layer 3 (32-bit , IPv4 128-bit IPv6)
  - Logical (IP) addresses are for universal communications that are independent of underlying physical networks
- Port address Layer 4 (16-bit)
  - Port addresses differentiate different processes
- Application-specific address Layer 7
  - Some applications have user-friendly addresses that are designed for that specific application, such as email address, URL.

# App-layer protocol defines

- **types of messages exchanged,**
  - e.g., request, response
- **message syntax:**
  - what fields in messages & how fields are delineated
- **message semantics**
  - meaning of information in fields
- **rules** for when and how processes send & respond to messages

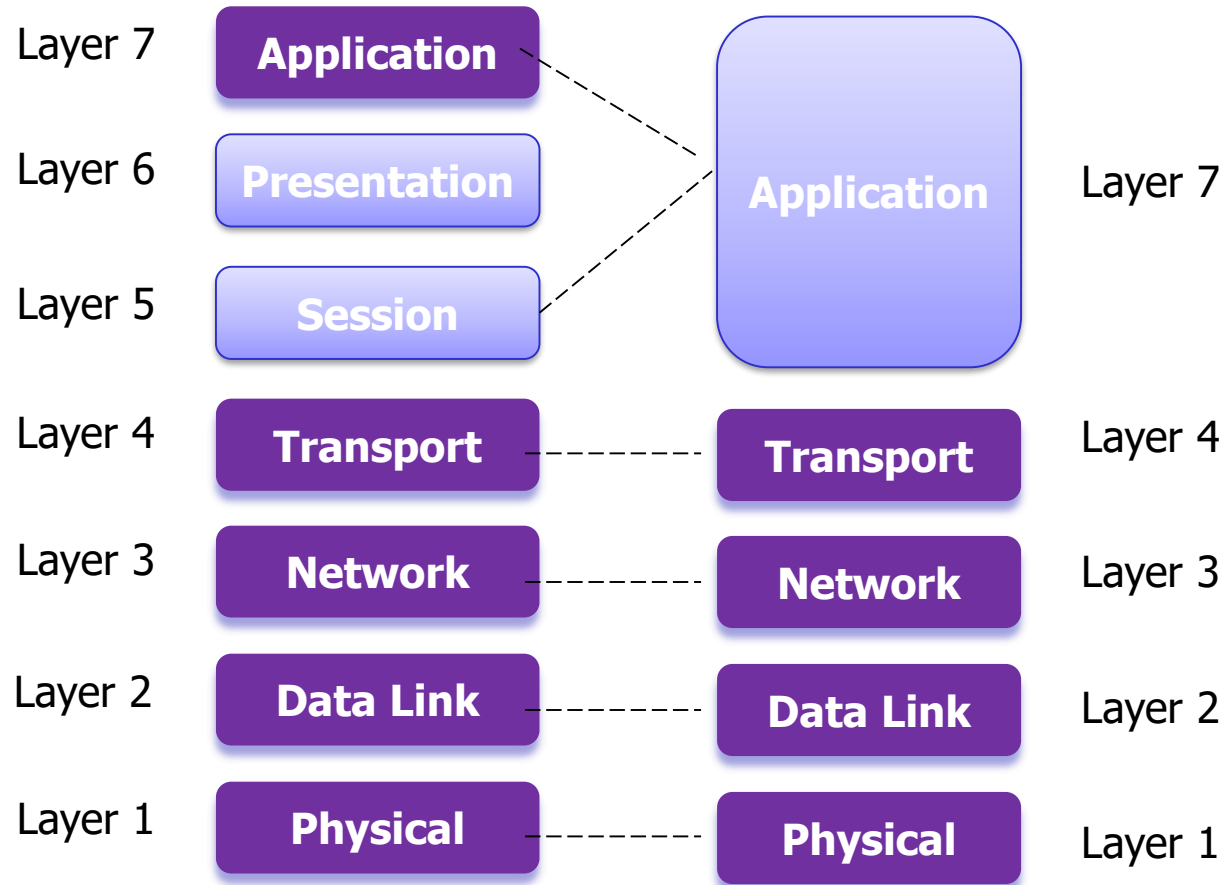
## **open protocols:**

- defined in RFCs (request for comments)
- allows for interoperability
- e.g., HTTP (web), SMTP (email)

## **proprietary protocols:**

- e.g., Skype

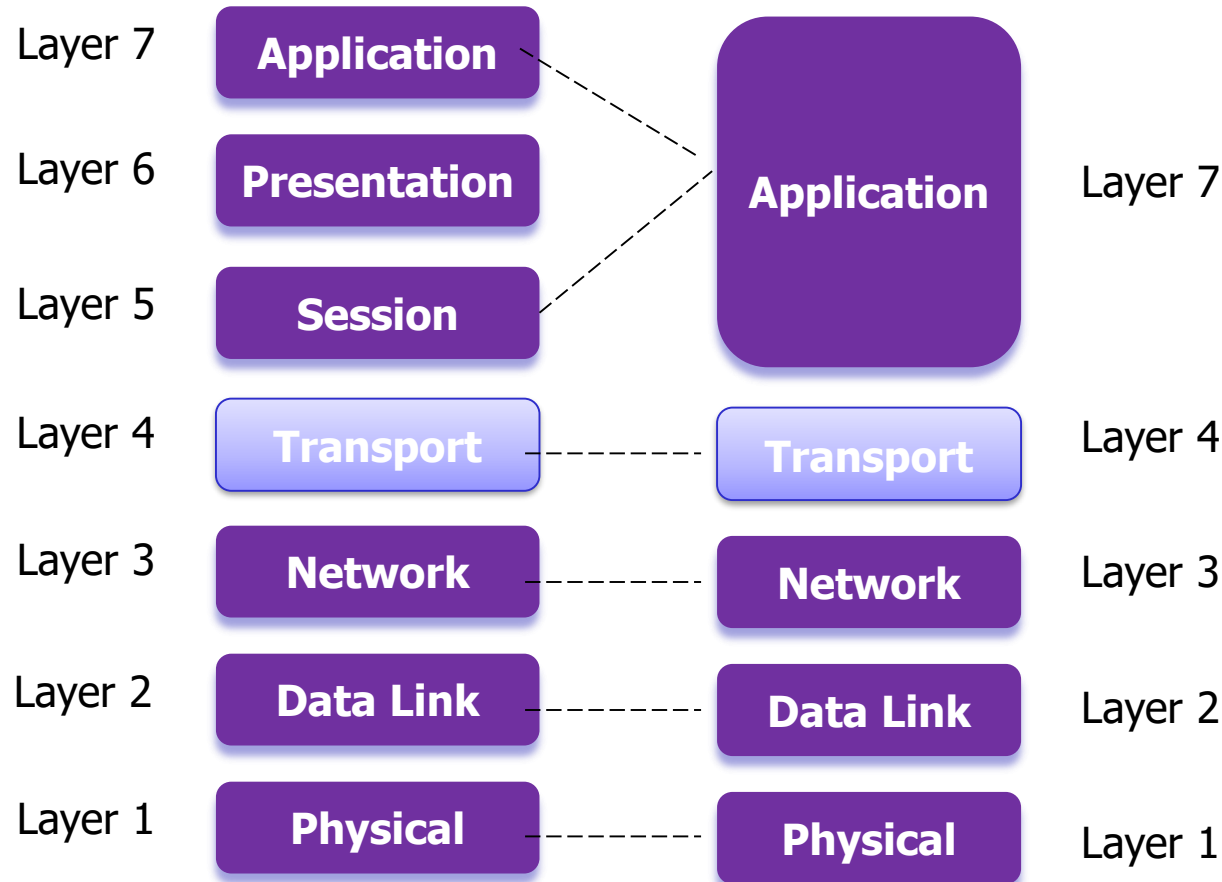
# Session + Presentation layers



# Session and Presentation Layers

- Theoretical only – not implemented in the Internet
- Session Layer (layer 5):
  - Takes care of a connection between two hosts for the lifetime of that connection
  - Authentication + authorisation (who is the user, what can they do)
  - In working Internet this is at the application layer.
- Presentation Layer (layer 6):
  - Translates data for an application.
  - For example takes care of detail of character set used to encode string of characters.
  - In working Internet this is at the application layer.

# Transport Layer



# Transport layer: outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

Part I (this week)

3.5 connection-oriented transport: TCP

- segment structure
- reliable data transfer
- flow control
- connection management

3.6 principles of congestion control

3.7 TCP congestion control

Part II (next week)

Numbers relate to chapter numbers in  
Course text: Kurose + Ross

# Transport Layer: outline

## 3.1 transport-layer services

## 3.2 multiplexing and demultiplexing

## 3.3 connectionless transport: UDP

## 3.4 principles of reliable data transfer

## 3.5 connection-oriented transport: TCP

- segment structure
- reliable data transfer
- flow control
- connection management

## 3.6 principles of congestion control

## 3.7 TCP congestion control

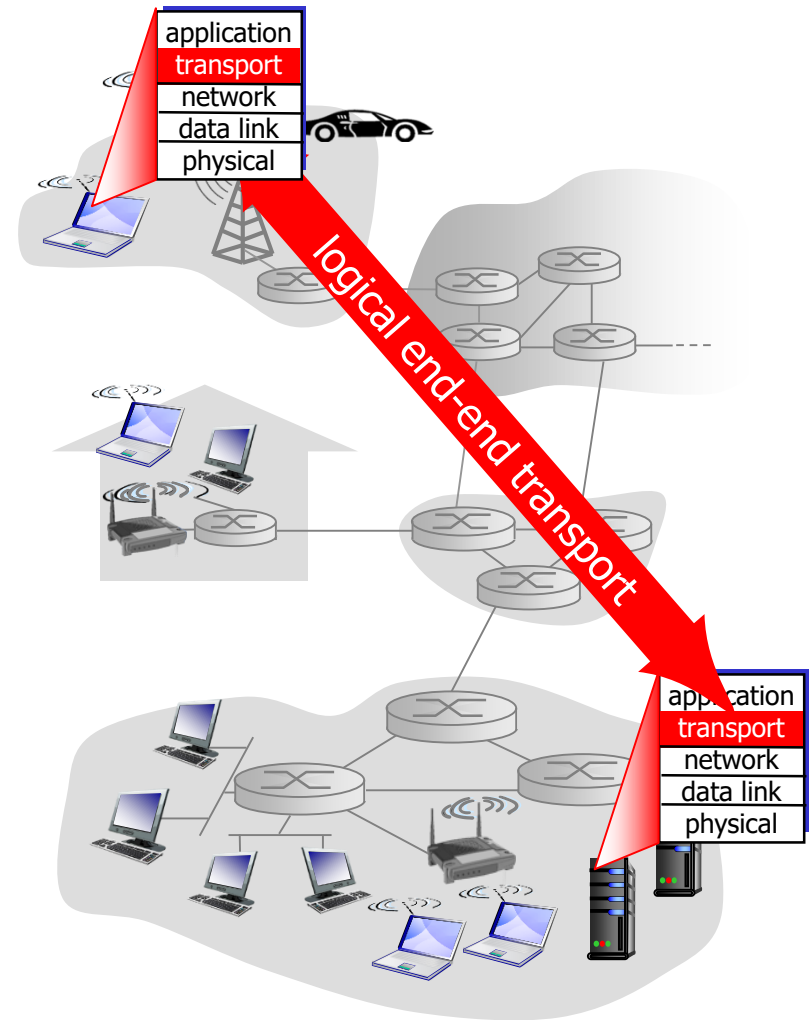


# How to remember the layers

- Please Do Not Throw Pizza + Sausage Away
- **P**lease – **P**hysical
- **D**o – **D**atalink
- **N**ot – **N**etwork
- **T**hrow – **T**ransport
- **S**ausage – **S**ession
- **P**izza – **P**resentation
- **A**way – **A**pplication

# Transport services and protocols

- provide *logical communication* between app processes running on different hosts
- transport protocols run in end systems
  - send side: breaks app messages into *segments*, passes to network layer
  - rcv side: reassembles segments into messages, passes to app layer
- more than one transport protocol available to apps
  - Internet: TCP and UDP



# Transport vs. network layer

- *network layer*: logical communication between hosts
- *transport layer*: logical communication between processes
  - relies on, enhances, network layer services

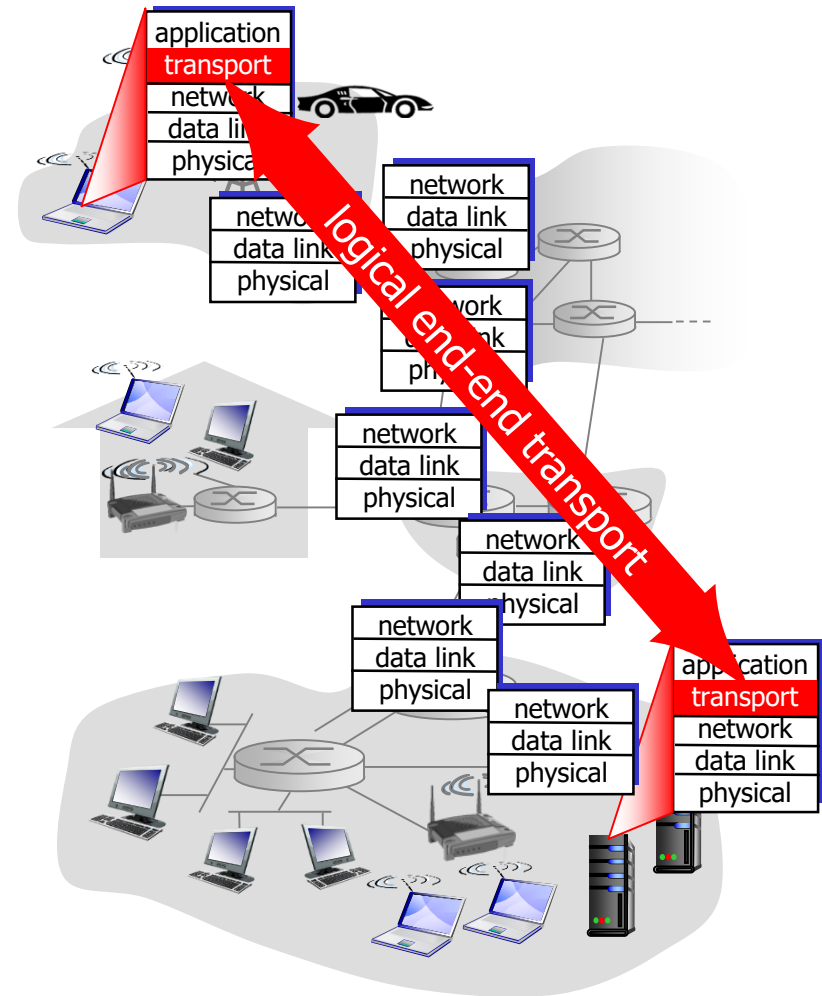
## *household analogy:*

*12 kids in Ann's house sending letters to 12 children in Bill's house:*

- hosts = houses
- processes = children
- app messages = letters in envelopes
- transport protocol = Ann and Bill who give letter to correct child
- network-layer protocol = postal service

# Internet transport-layer protocols

- reliable, in-order delivery (TCP)
  - congestion control
  - flow control
  - connection setup
- unreliable, unordered delivery: UDP
  - no-frills extension of “best-effort” IP
- services not available:
  - delay guarantees
  - bandwidth guarantees



# Transport Layer: outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

- segment structure
- reliable data transfer
- flow control
- connection management

3.6 principles of congestion control

3.7 TCP congestion control

# Multiplexing/demultiplexing

## ■ Multiplexing (Mux):

- Combining several streams of data into a single stream.
- Example – your phone is browsing the web, refreshing your email, connecting to wechat at the same time. All these connections are sent over the same link.

## ■ Demultiplexing (Demux)

- The opposite process. A stream of data is separated out into its individual components.
- The stream of packets the phone received is split up and sent to the appropriate program for web, email, wechat.

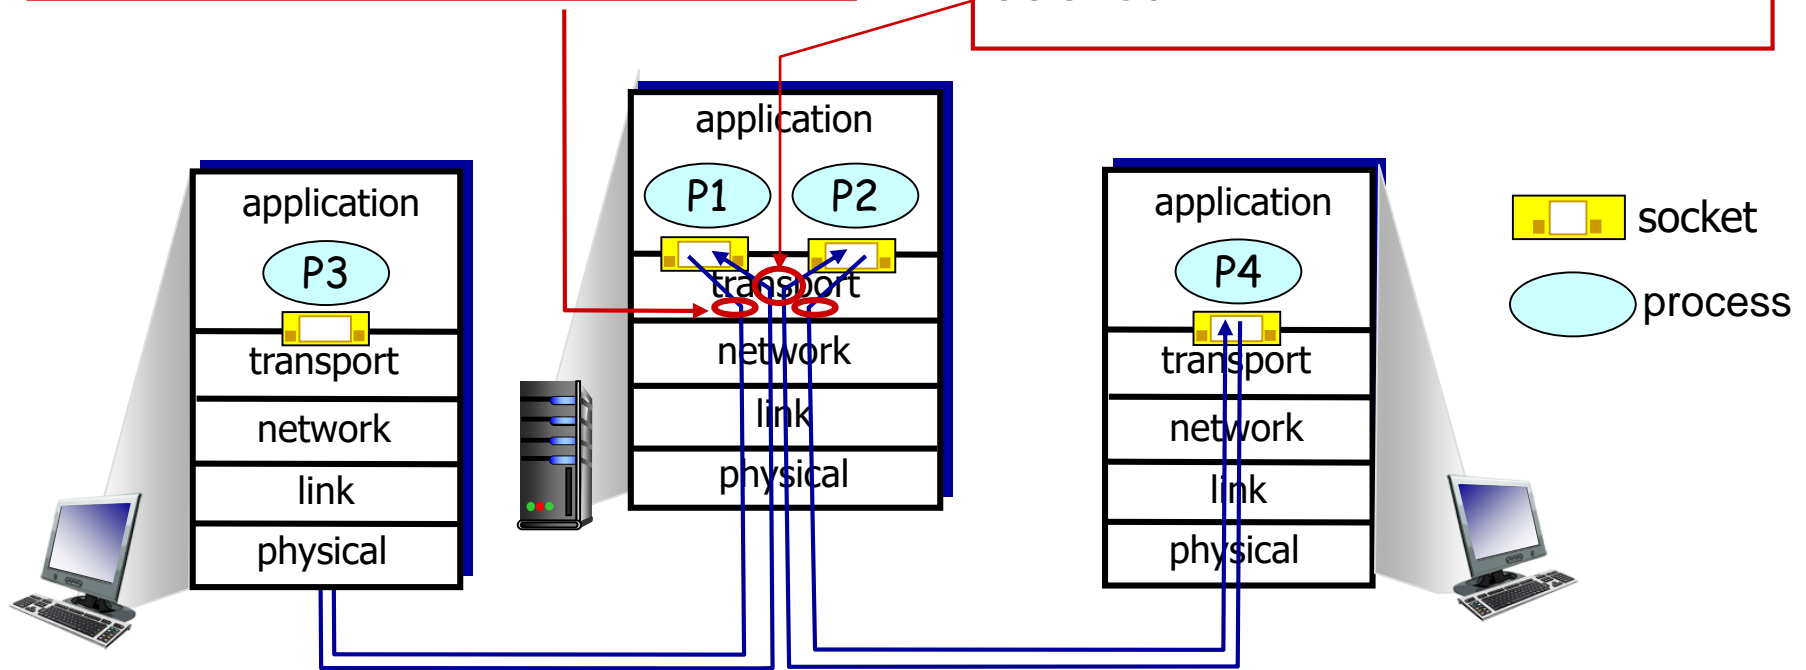
# Multiplexing/demultiplexing

## *multiplexing at sender:*

handle data from multiple sockets, add transport header (later used for demultiplexing)

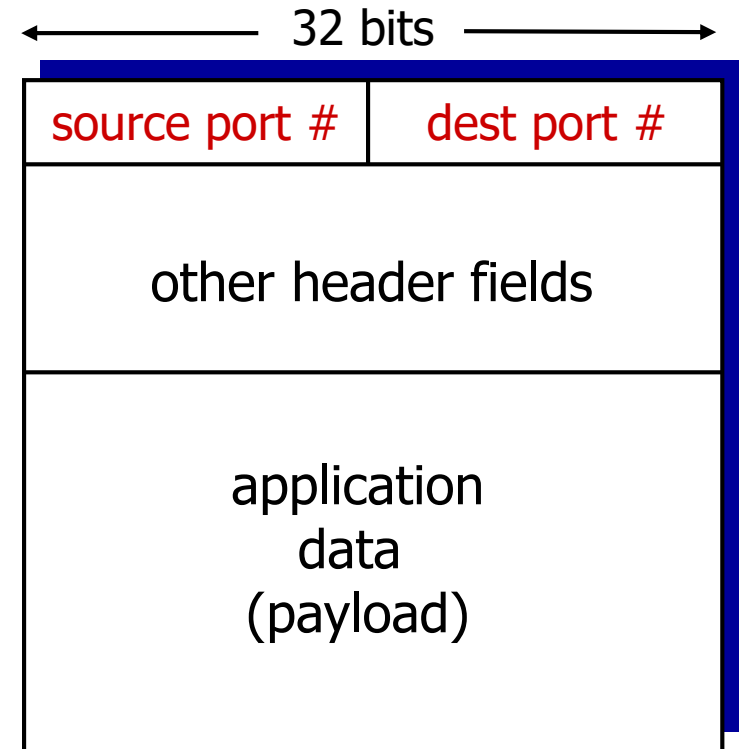
## *demultiplexing at receiver:*

use header info to deliver received segments to correct socket



# How demultiplexing works

- host receives IP datagrams
  - each datagram has source IP address, destination IP address
  - each datagram carries one transport-layer segment
  - each segment has source, destination port number
- host uses *IP addresses & port numbers* to direct segment to appropriate socket



TCP/UDP segment format



# Connectionless demultiplexing

- *recall*: created socket has host-local port #:

```
DatagramSocket mySocket1  
= new DatagramSocket(12534) ;
```

- *recall*: when creating datagram to send into UDP socket, must specify

- destination IP address
- destination port #

- 
- when host receives UDP segment:

- checks destination port # in segment
- directs UDP segment to socket with that port #



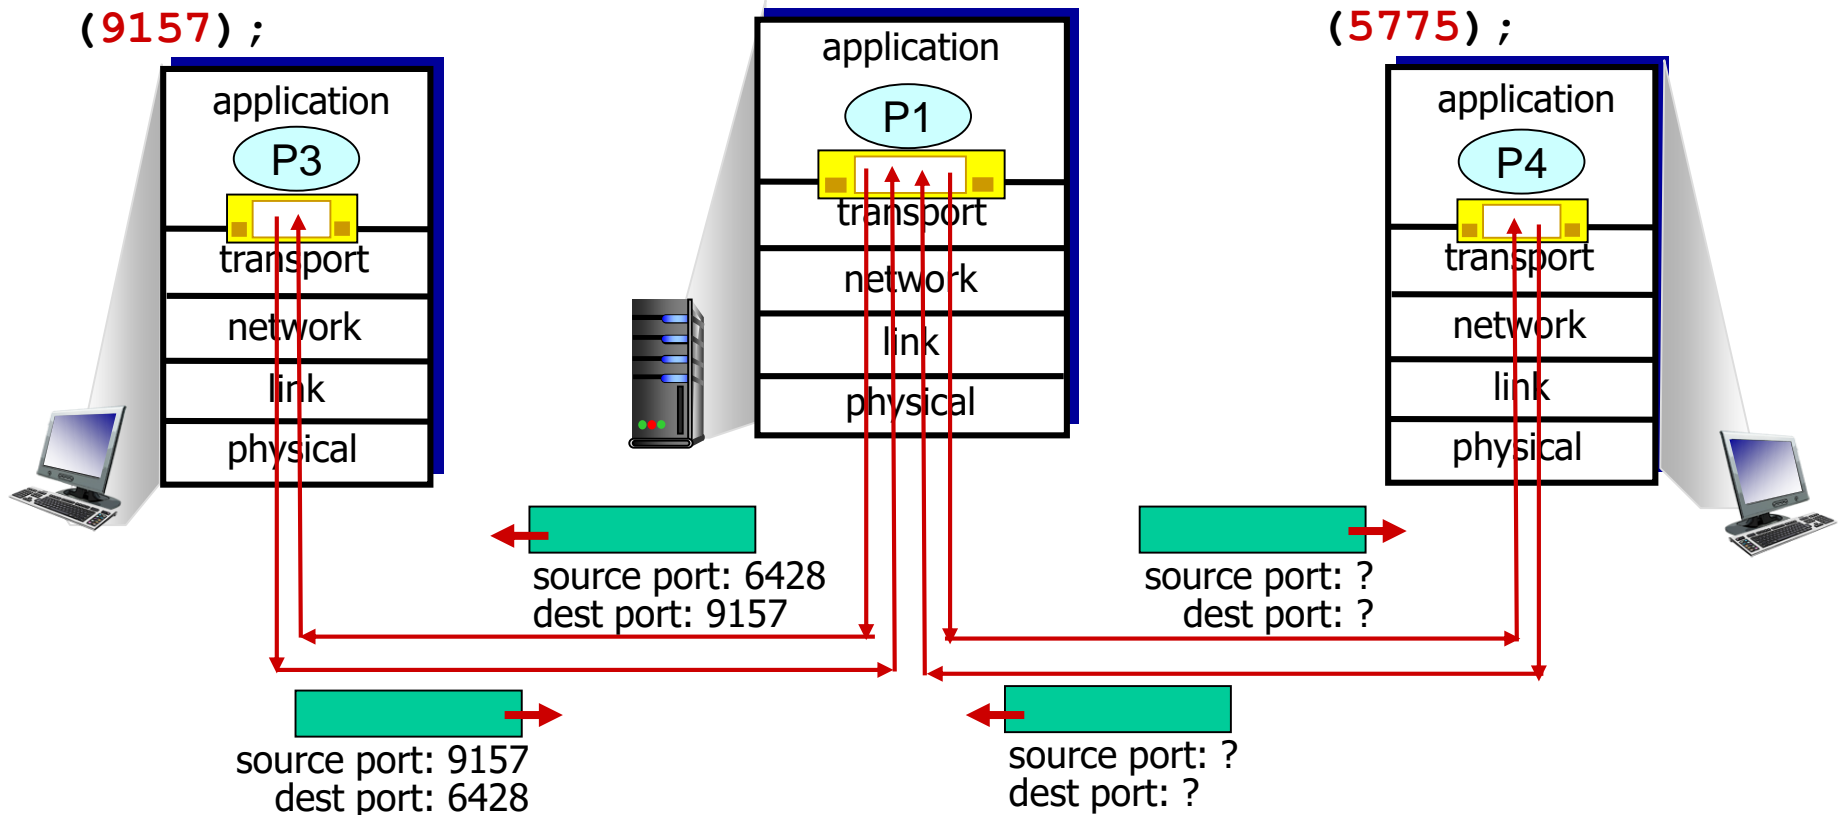
IP datagrams with *same dest. port #*, but different source IP addresses and/or source port numbers will be directed to *same socket* at dest

# Connectionless demux: example

```
DatagramSocket  
mySocket2 = new  
DatagramSocket  
(9157);
```

```
DatagramSocket  
serverSocket = new  
DatagramSocket  
(6428);
```

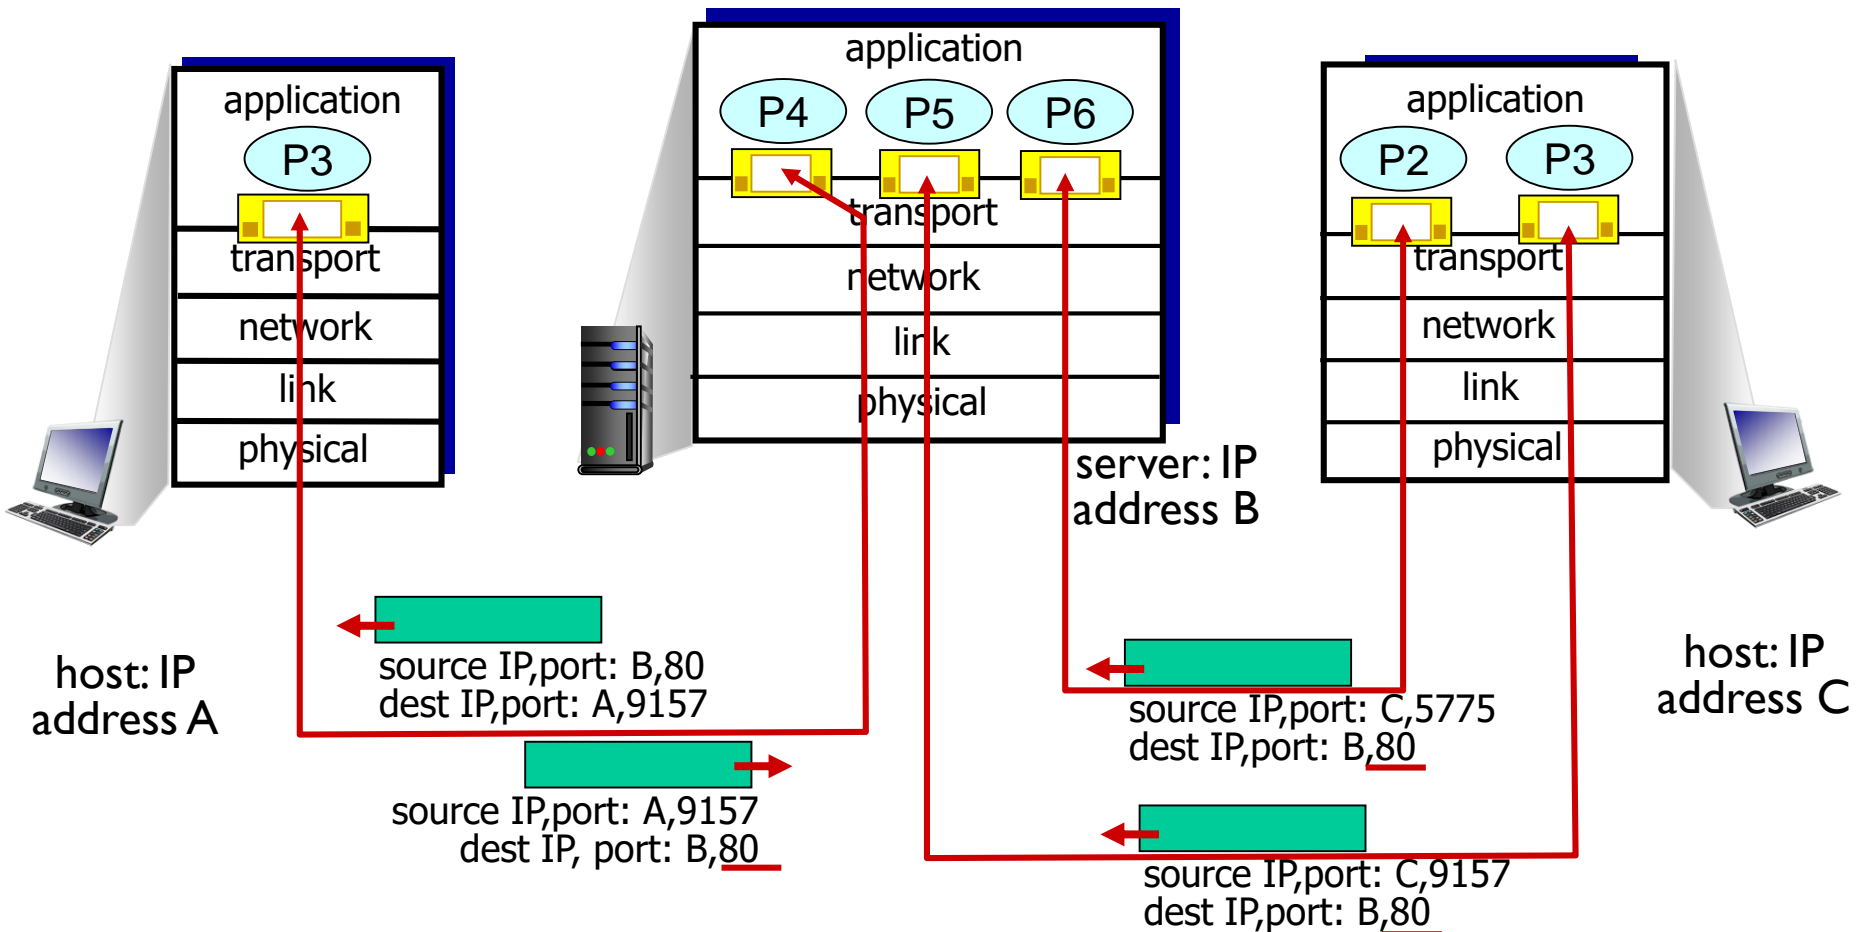
```
DatagramSocket  
mySocket1 = new  
DatagramSocket  
(5775);
```



# Connection-oriented demux

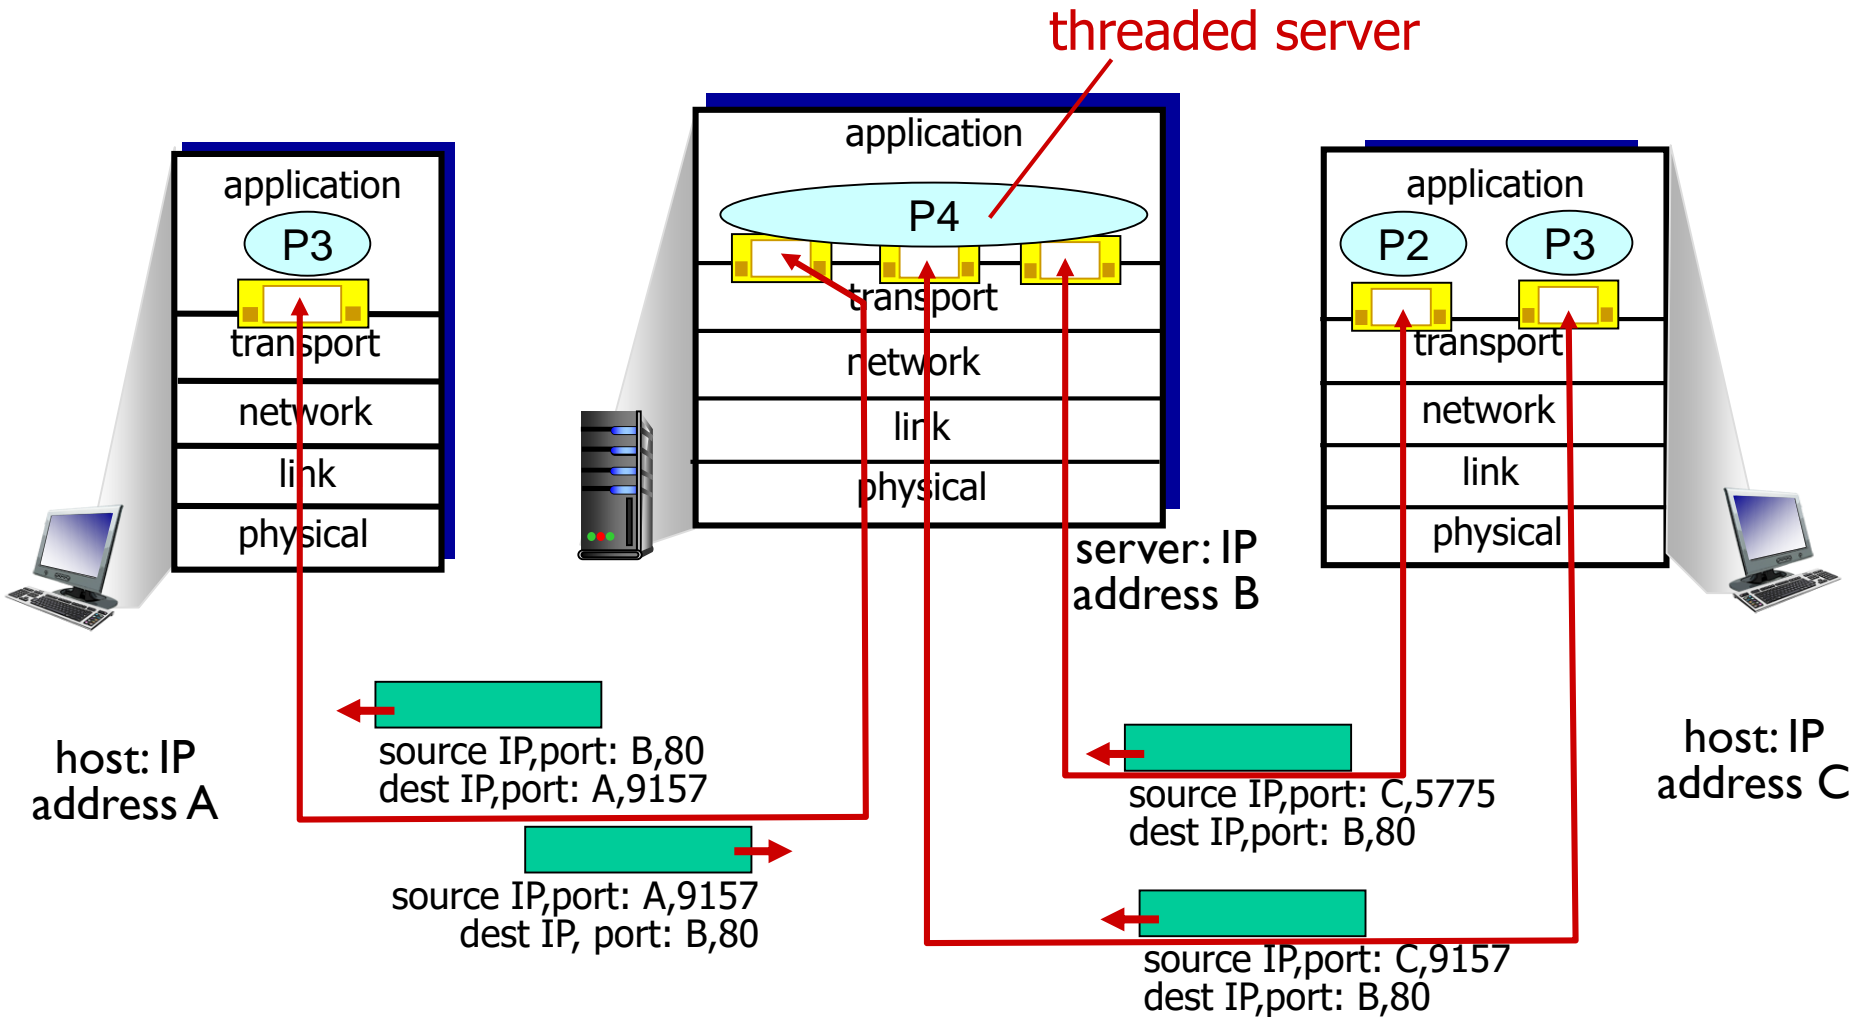
- TCP socket identified by 4-tuple:
  - source IP address
  - source port number
  - dest IP address
  - dest port number
- demux: receiver uses all four values to direct segment to appropriate socket
- server host may support many simultaneous TCP sockets:
  - each socket identified by its own 4-tuple
- web servers have different sockets for each connecting client
  - non-persistent HTTP will have different socket for each request

# Connection-oriented demux: example



three segments, all destined to IP address: B,  
dest port: 80 are demultiplexed to *different* sockets

# Connection-oriented demux: example



# Typical port numbers for applications

- Port 80: Standard HTTP (Hypertext Transfer Protocol) – browse the web
- Port 22: ssh (secure shell) log in from remote computer
- Port 25: SMTP (simple mail transfer protocol) – send mail via this host
- Port 143: IMAP (Internet Message Access Protocol) – read your email
- Port 443: HTTPS (secure HTTP) browse the web securely
- These are just conventions, you could change them if you so wished on your own host.

# Transport layer outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

- segment structure
- reliable data transfer
- flow control
- connection management

3.6 principles of congestion control

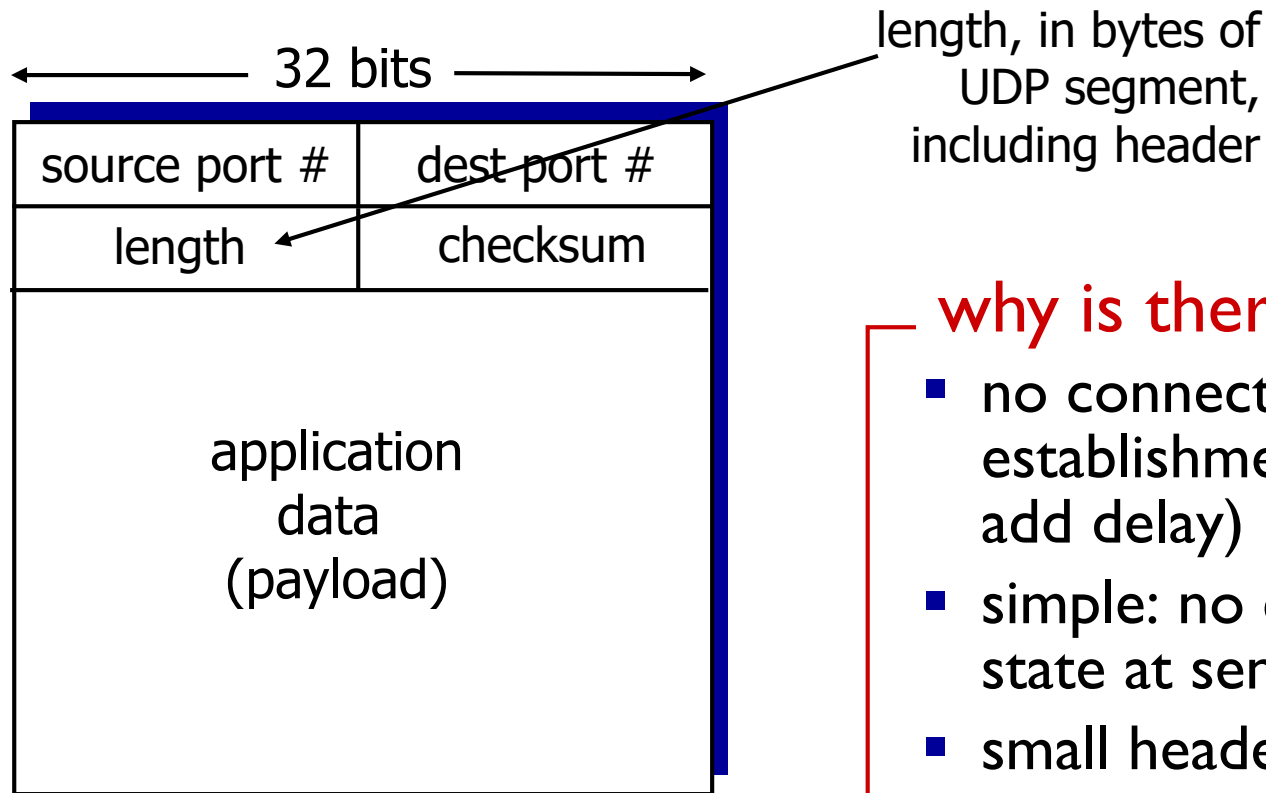
3.7 TCP congestion control

# UDP: User Datagram Protocol [RFC 768]

- Simplest usable Internet transport protocol
- “best effort” service, UDP segments may be:
  - lost
  - delivered out-of-order to app
- *connectionless*:
  - no “handshaking” between UDP sender, receiver (can send immediately without asking first)
  - each UDP packet handled independently of others
- UDP use:
  - streaming multimedia apps (loss tolerant, rate sensitive)
  - DNS
  - SNMP
- reliable transfer over UDP:
  - add reliability at application layer
  - application-specific error recovery!



# UDP: segment header



UDP segment format

## why is there a UDP?

- no connection establishment (which can add delay)
- simple: no connection state at sender, receiver
- small header size
- no congestion control: UDP can blast away as fast as desired

# UDP checksum

*Goal:* detect “errors” (e.g., flipped bits – 0 becomes 1 or 1 becomes 0) in transmitted segment

## sender:

- treat segment contents, including header fields, as sequence of 16-bit integers
- checksum: addition (one's complement sum) of segment contents
- sender puts checksum value into UDP checksum field

## receiver:

- compute checksum of received segment
- check if computed checksum equals checksum field value:
  - NO - error detected
  - YES - no error detected.  
*But maybe errors nonetheless? More later*  
....

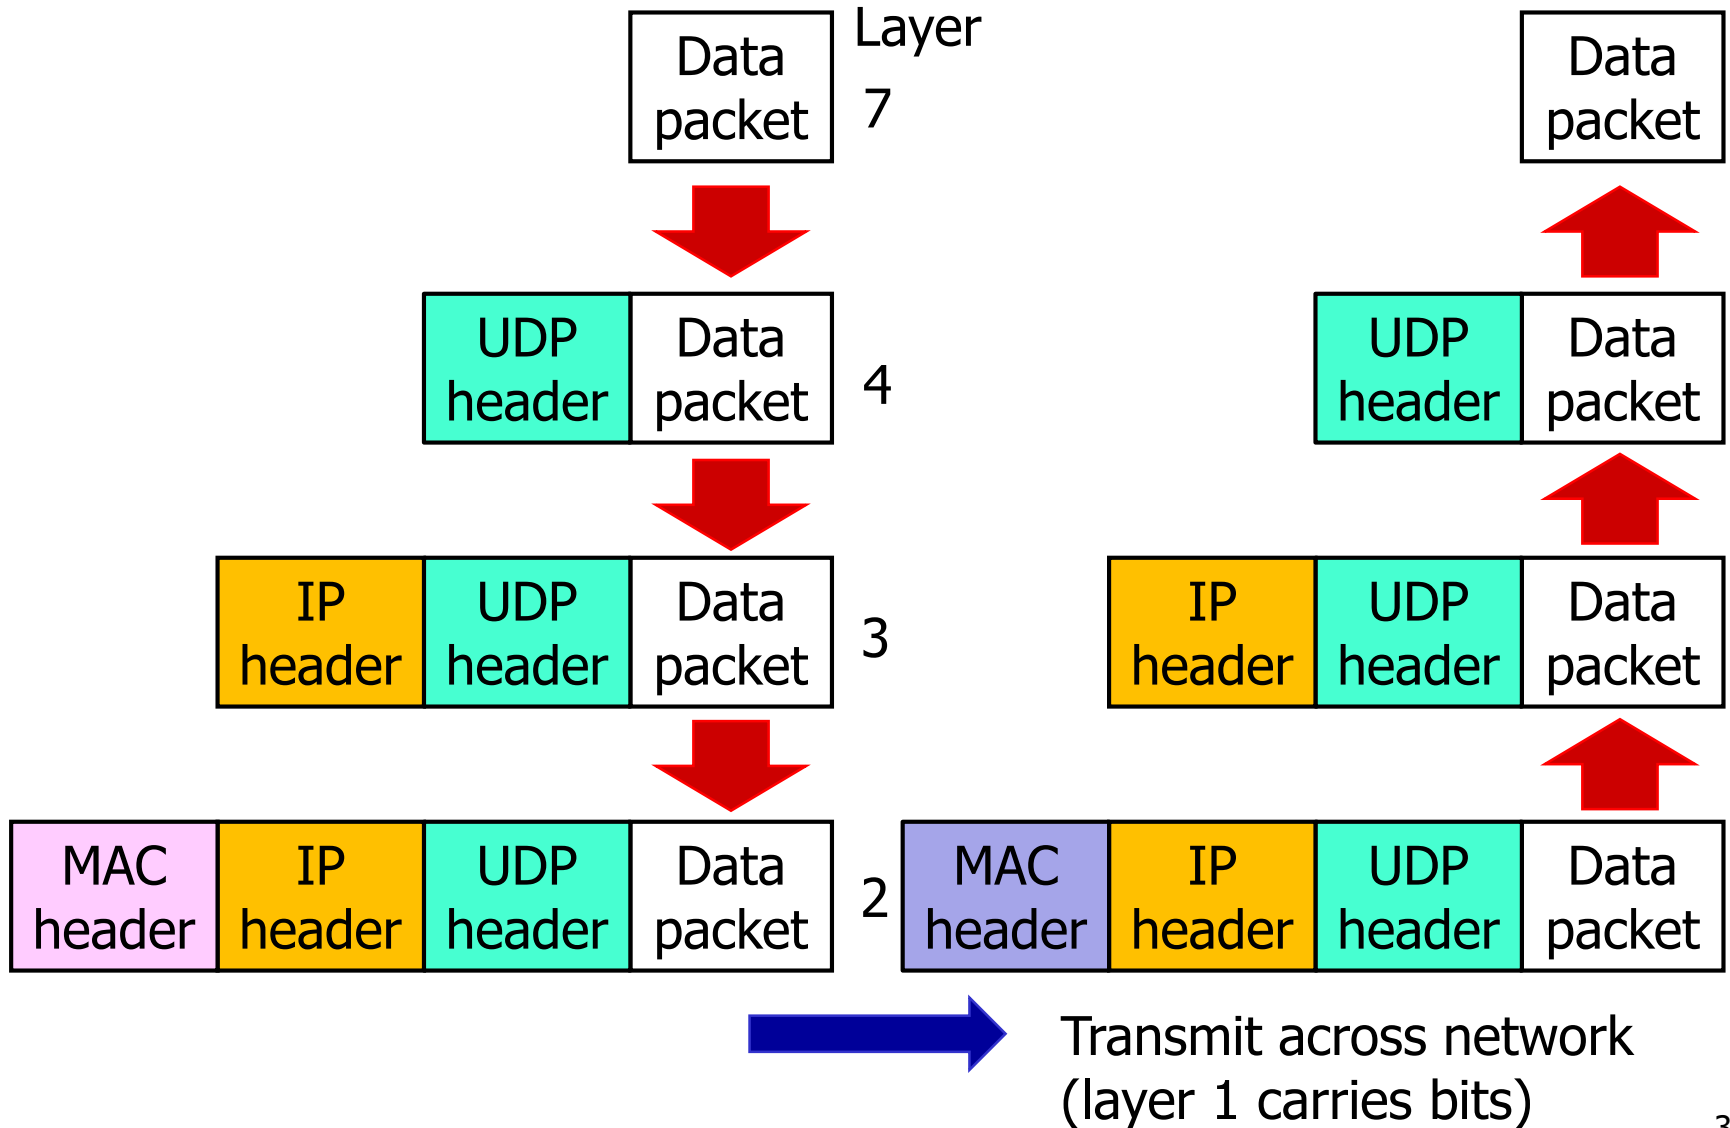
# Internet checksum: example

example: add two 16-bit integers

	1	1	1	0	0	1	1	0	0	1	1	0	0	1	1	0
	1	1	0	1	0	1	0	1	0	1	0	1	0	1	0	1
<hr/>																
wraparound	1	1	0	1	1	1	0	1	1	1	0	1	1	1	0	1
<hr/>																
sum	1	0	1	1	1	0	1	1	1	0	1	1	1	1	0	0
checksum	0	1	0	0	0	1	0	0	0	1	0	0	0	0	1	1

*Note:* when adding numbers, a carryout from the most significant bit needs to be added to the result

# UDP Encapsulation/decapsulation



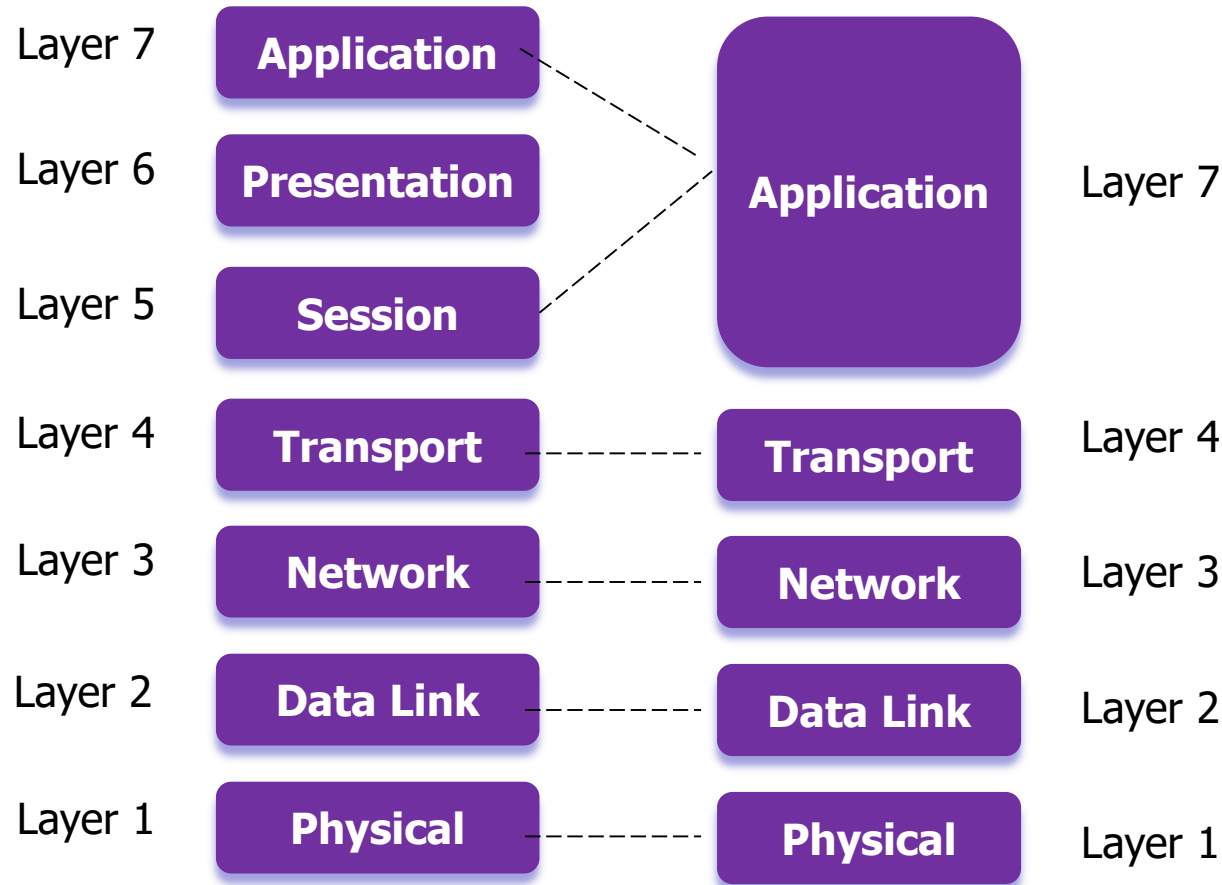
# What have we learned?

- Brief look at Application Layer (layer seven) protocols which programmers use.
- Transport Layer introduction (layer four).
  - Deliver data to application sockets themselves.
  - Responsible for “end-to-end” connection.
- User Datagram Protocol (transport layer):
  - Simple, connectionless.
  - Packets may be lost and reordered.
  - Useful for simple situations or real-time.

# Structure of course

- Week 1 (8<sup>th</sup>-12<sup>th</sup> October)
  - Introduction to IP Networks
  - The Transport layer (part I)
- Week 2 (29<sup>th</sup> October - 2<sup>nd</sup> November)
  - The Transport layer (part II)
  - The Network layer (part I)
  - Class test
- Week 3 (19<sup>th</sup>-23<sup>rd</sup> November)
  - The Network layer (part II)
  - The Data link layer (part I)
  - Router lab tutorial (assessed lab work after this week)
- Week 4 (10<sup>th</sup>-14<sup>th</sup> December)
  - The Data link layer (part II)
  - Network management and security
  - Class test

# ISO/OSI (left) vs TCP/IP (right)

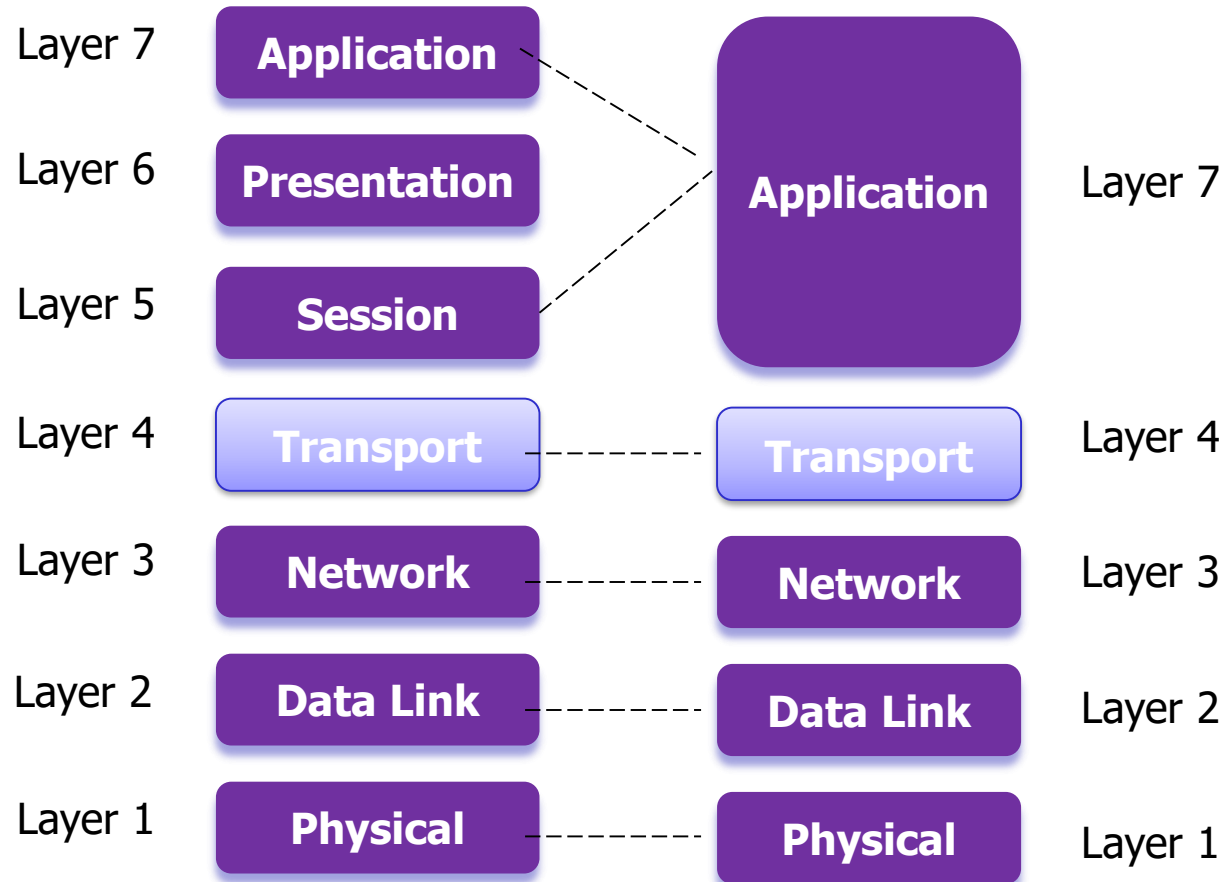


# How to remember the layers

- Please Do Not Throw Pizza + Sausage Away
- **P**lease – **P**hysical
- **D**o – **D**atalink
- **N**ot – **N**etwork
- **T**hrow – **T**ransport
- **S**ausage – **S**ession
- **P**izza – **P**resentation
- **A**way – **A**pplication



# Transport Layer



# Week 1: Transport Layer (part II)

## Our goals:

- understand principles behind transport layer services:
  - multiplexing, demultiplexing
  - reliable data transfer
  - flow control
  - congestion control
- learn about Internet transport layer protocols:
  - UDP: connectionless transport
  - TCP: connection-oriented reliable transport
  - TCP congestion control

# Transport layer outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

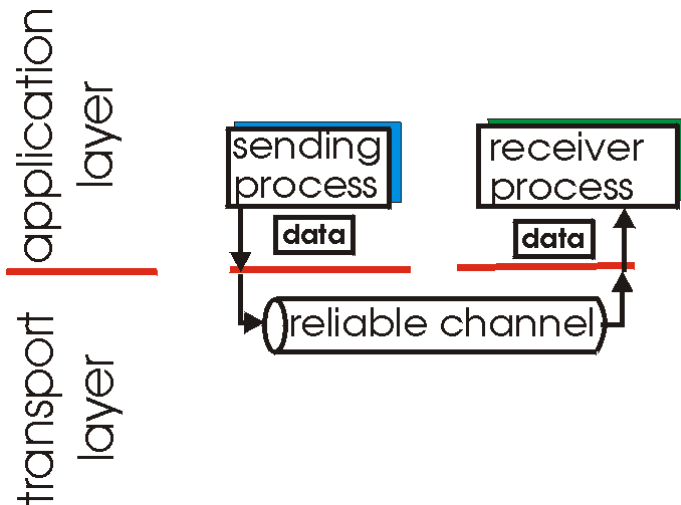
- segment structure
- reliable data transfer
- flow control
- connection management

3.6 principles of congestion control

3.7 TCP congestion control

# Principles of reliable data transfer

- important in application, transport, link layers
  - top-10 list of important networking topics!

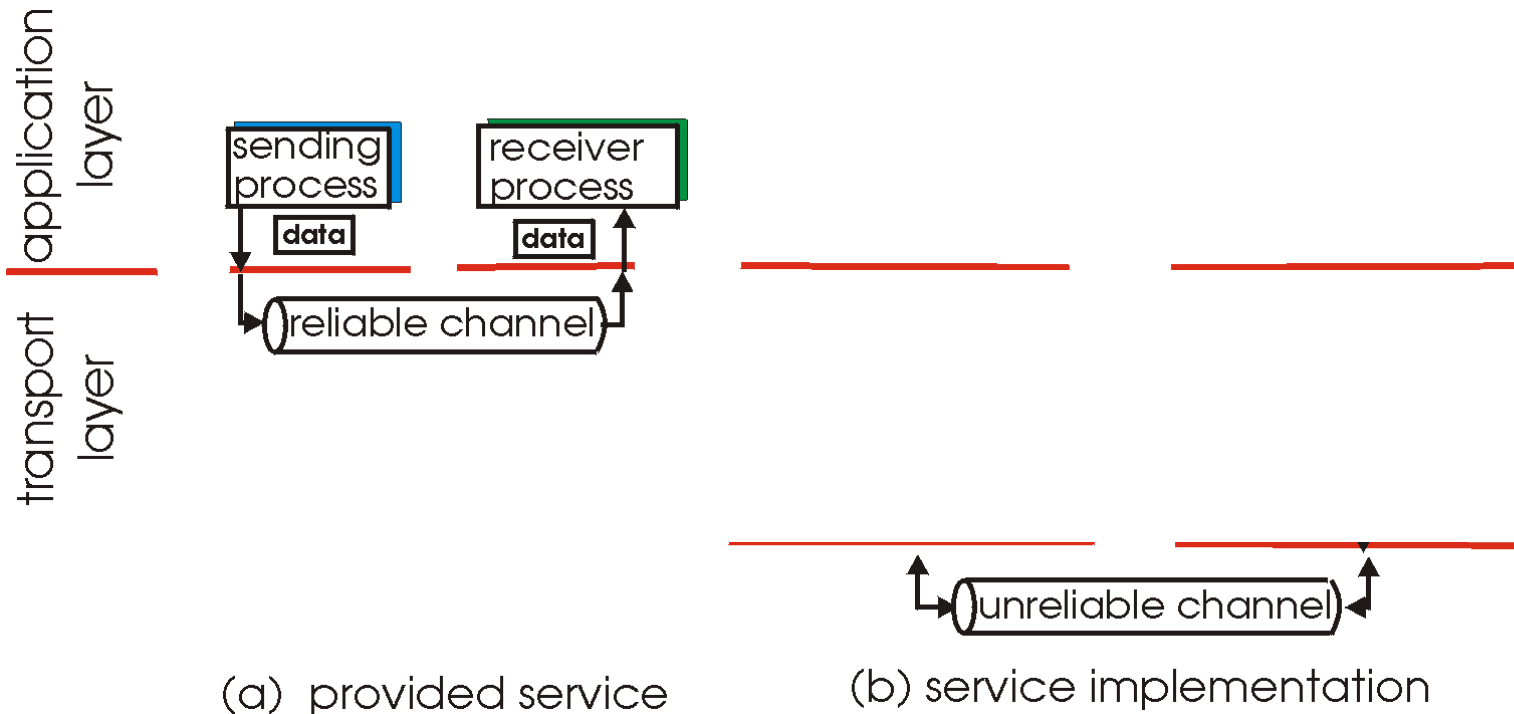


(a) provided service

- characteristics of unreliable channel will determine complexity of reliable data transfer protocol (rdt)

# Principles of reliable data transfer

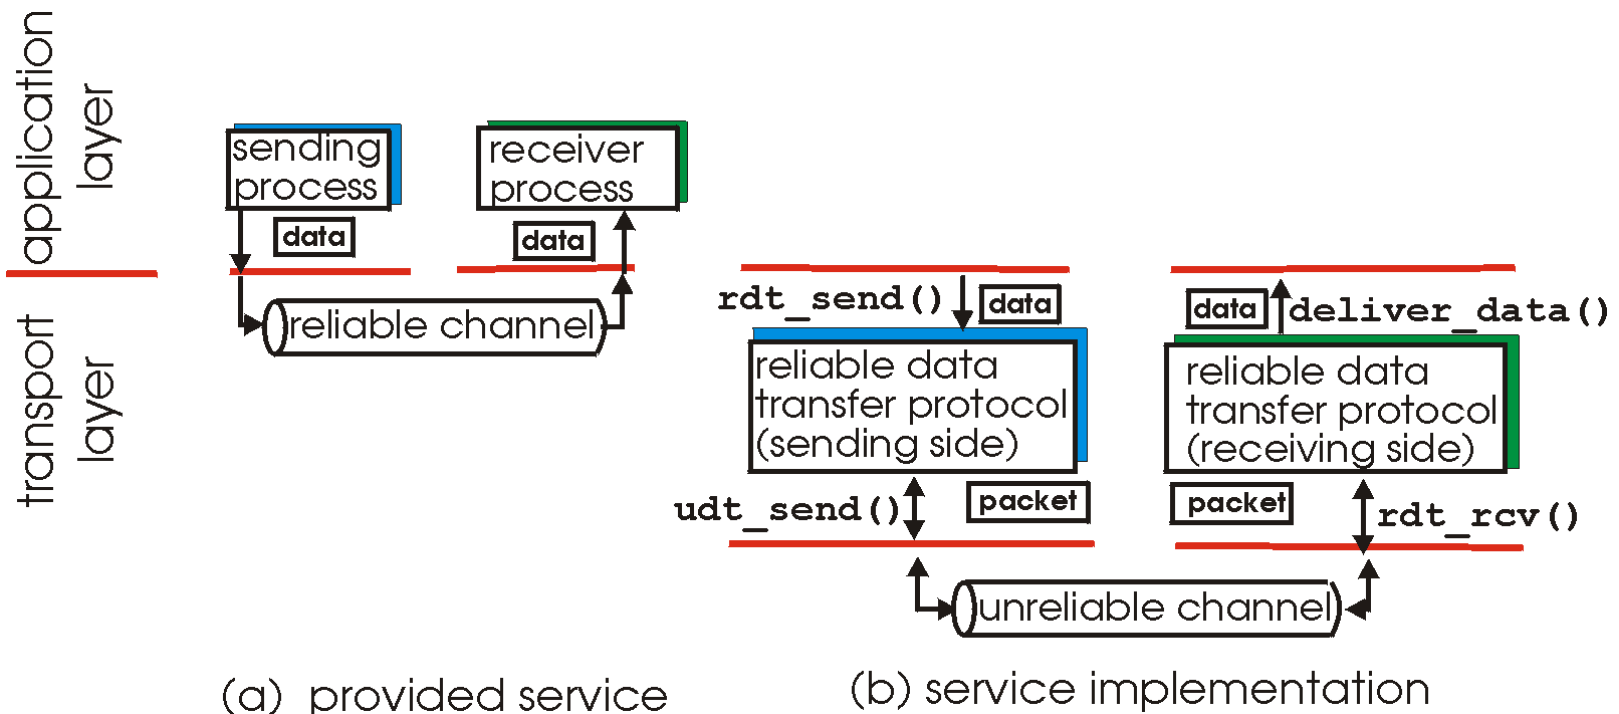
- important in application, transport, link layers
  - top-10 list of important networking topics!



- characteristics of unreliable channel will determine complexity of reliable data transfer protocol (rdt)

# Principles of reliable data transfer

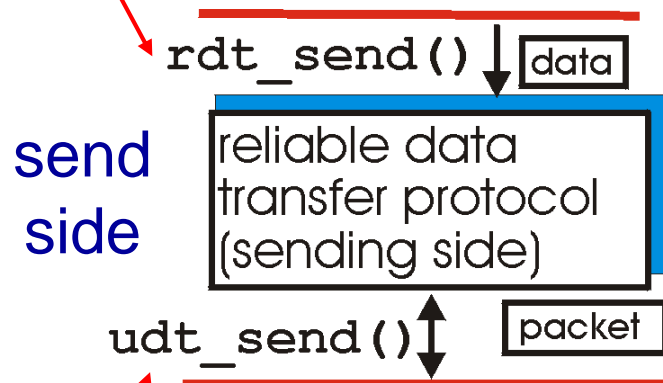
- important in application, transport, link layers
  - top-10 list of important networking topics!



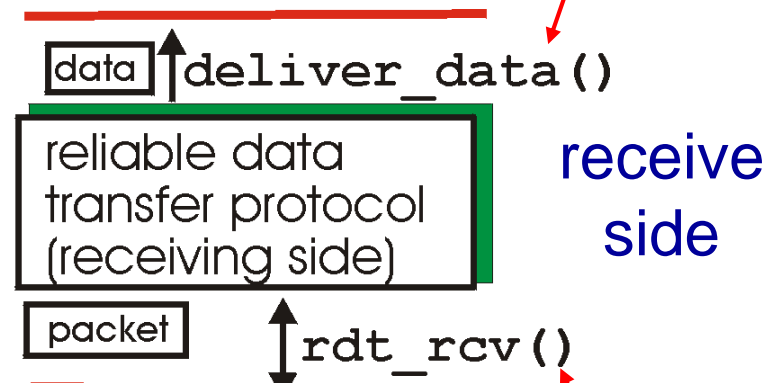
- characteristics of unreliable channel will determine complexity of reliable data transfer protocol (rdt)

# Reliable data transfer: getting started

**rdt\_send()** : called from above,  
(e.g., by app.). Passed data to  
deliver to receiver upper layer



**deliver\_data()** : called by  
**rdt** to deliver data to upper



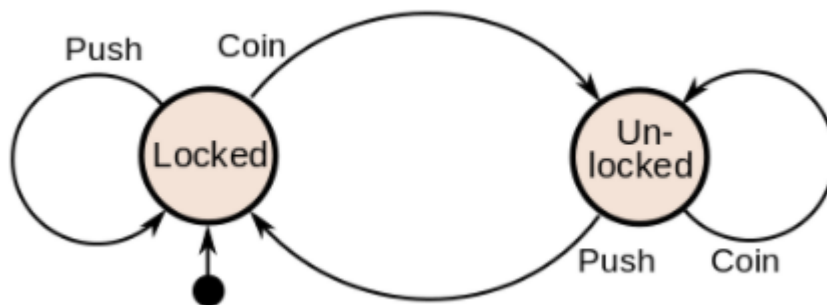
**udt\_send()** : called by rdt,  
to transfer packet over  
unreliable channel to receiver

**rdt\_rcv()** : called when packet  
arrives on rcv-side of channel



# What is a Finite State Machine

- A finite state machine shows how a system works with states and transitions between them.
- Consider a coin operated turnstile – we put in a coin, it allows us in.
- It has two states "locked" "unlocked".
- It has two events "push" and "insert coin"



Diagrams from wikipedia



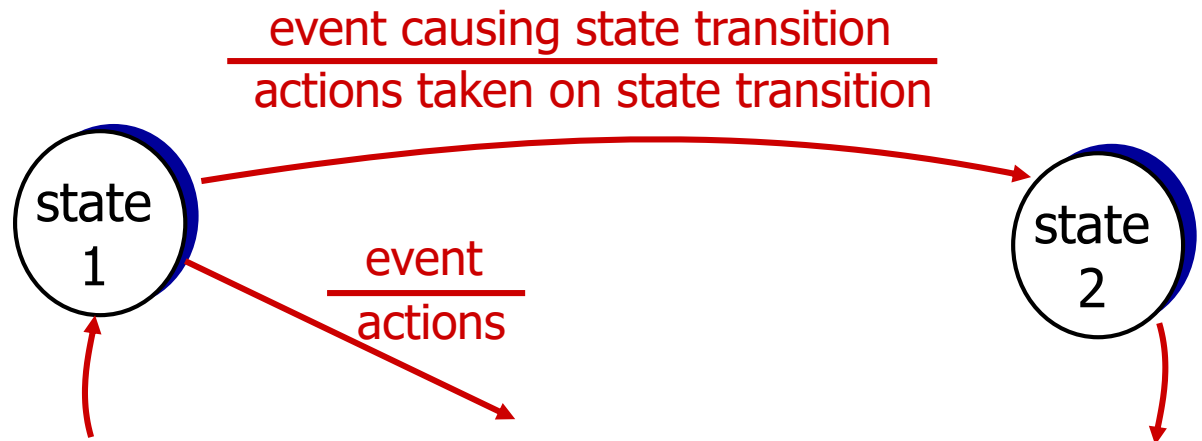


# Reliable data transfer: getting started

We'll:

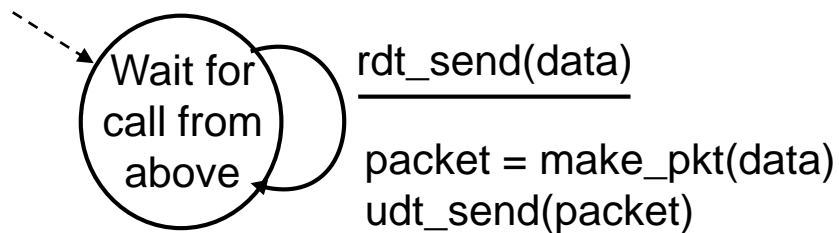
- incrementally develop sender, receiver sides of reliable data transfer protocol (rdt)
- consider only unidirectional data transfer
  - but control info will flow on both directions!
- use finite state machines (FSM) to specify sender, receiver

**state:** when in this “state” next state uniquely determined by next event

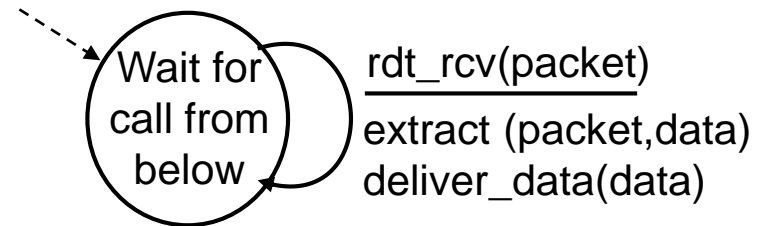


# rdt1.0: reliable transfer over a reliable channel

- underlying channel perfectly reliable
  - no bit errors
  - no loss of packets
- separate FSMs for sender, receiver:
  - sender sends data into underlying channel
  - receiver reads data from underlying channel



sender



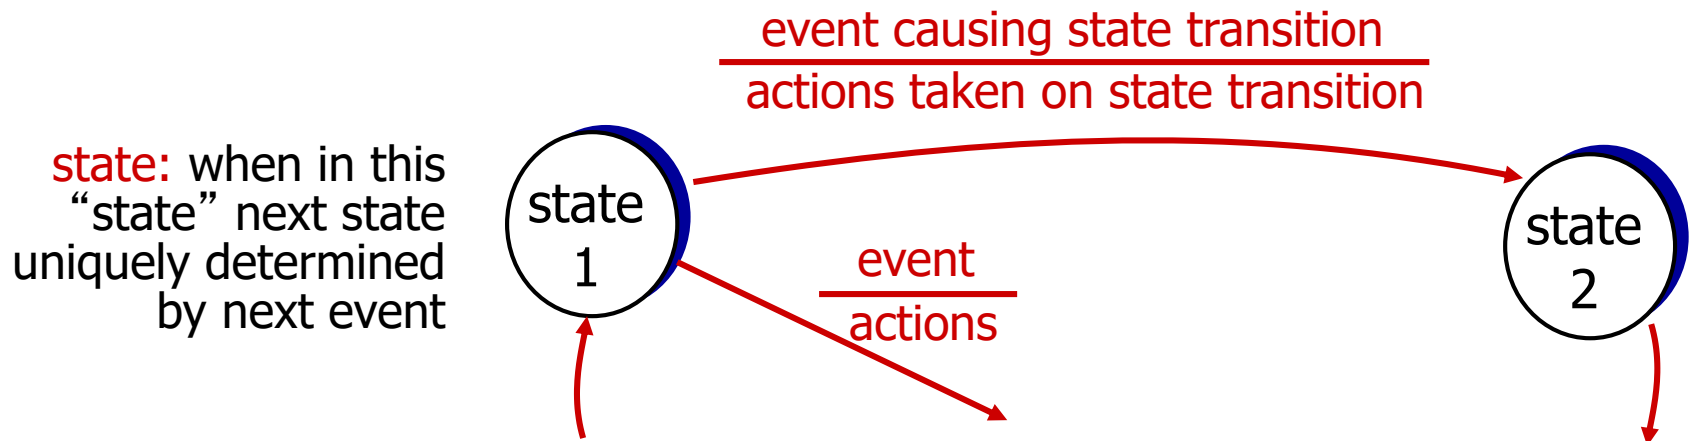
receiver

# rdt2.0: channel with bit errors

- We drop the assumption that the channel has no errors.
- Now the channel could have bit errors.
- We need to be able to detect and recover from errors in transmission.
- Consider this question:
  - How do “humans” recover from errors in conversation?

# Reliable data transfer: getting started

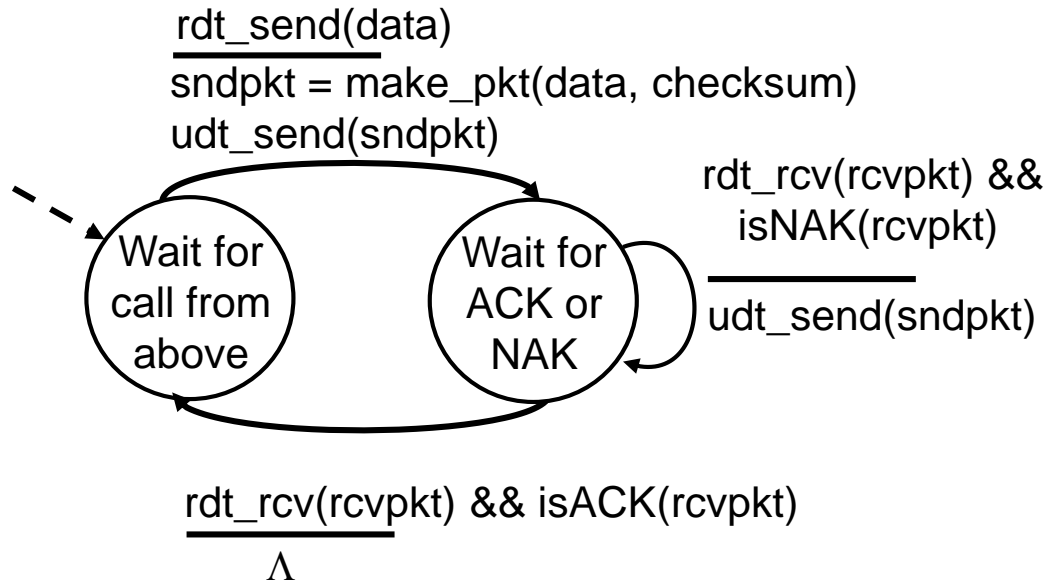
- incrementally develop sender, receiver sides of reliable data transfer protocol (rdt)
- consider only unidirectional data transfer
  - but control info will flow on both directions!
- use finite state machines (FSM) to specify sender, receiver



# rdt2.0: channel with bit errors

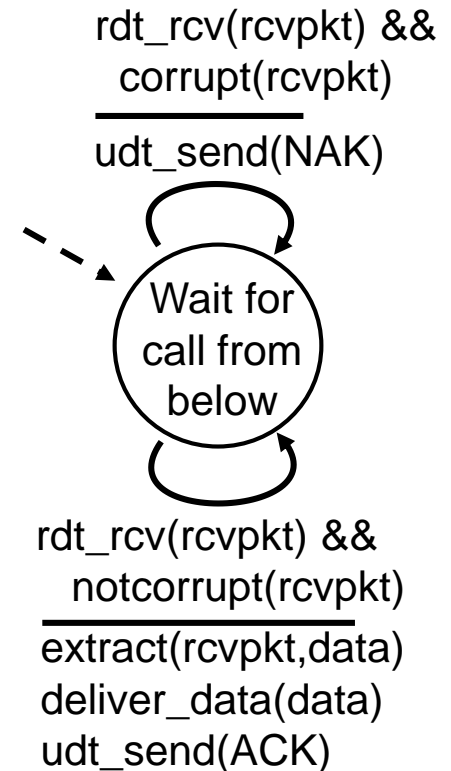
- underlying channel may flip bits in packet
  - checksum to detect bit errors
- *the question*: how to recover from errors:
  - *acknowledgements (ACKs)*: receiver explicitly tells sender that pkt received OK
  - *negative acknowledgements (NAKs)*: receiver explicitly tells sender that pkt had errors
  - sender retransmits pkt on receipt of NAK
- new mechanisms in `rdt2.0` (beyond `rdt1.0`):
  - error detection
  - feedback: control msgs (ACK,NAK) from receiver to sender

# rdt2.0: FSM specification

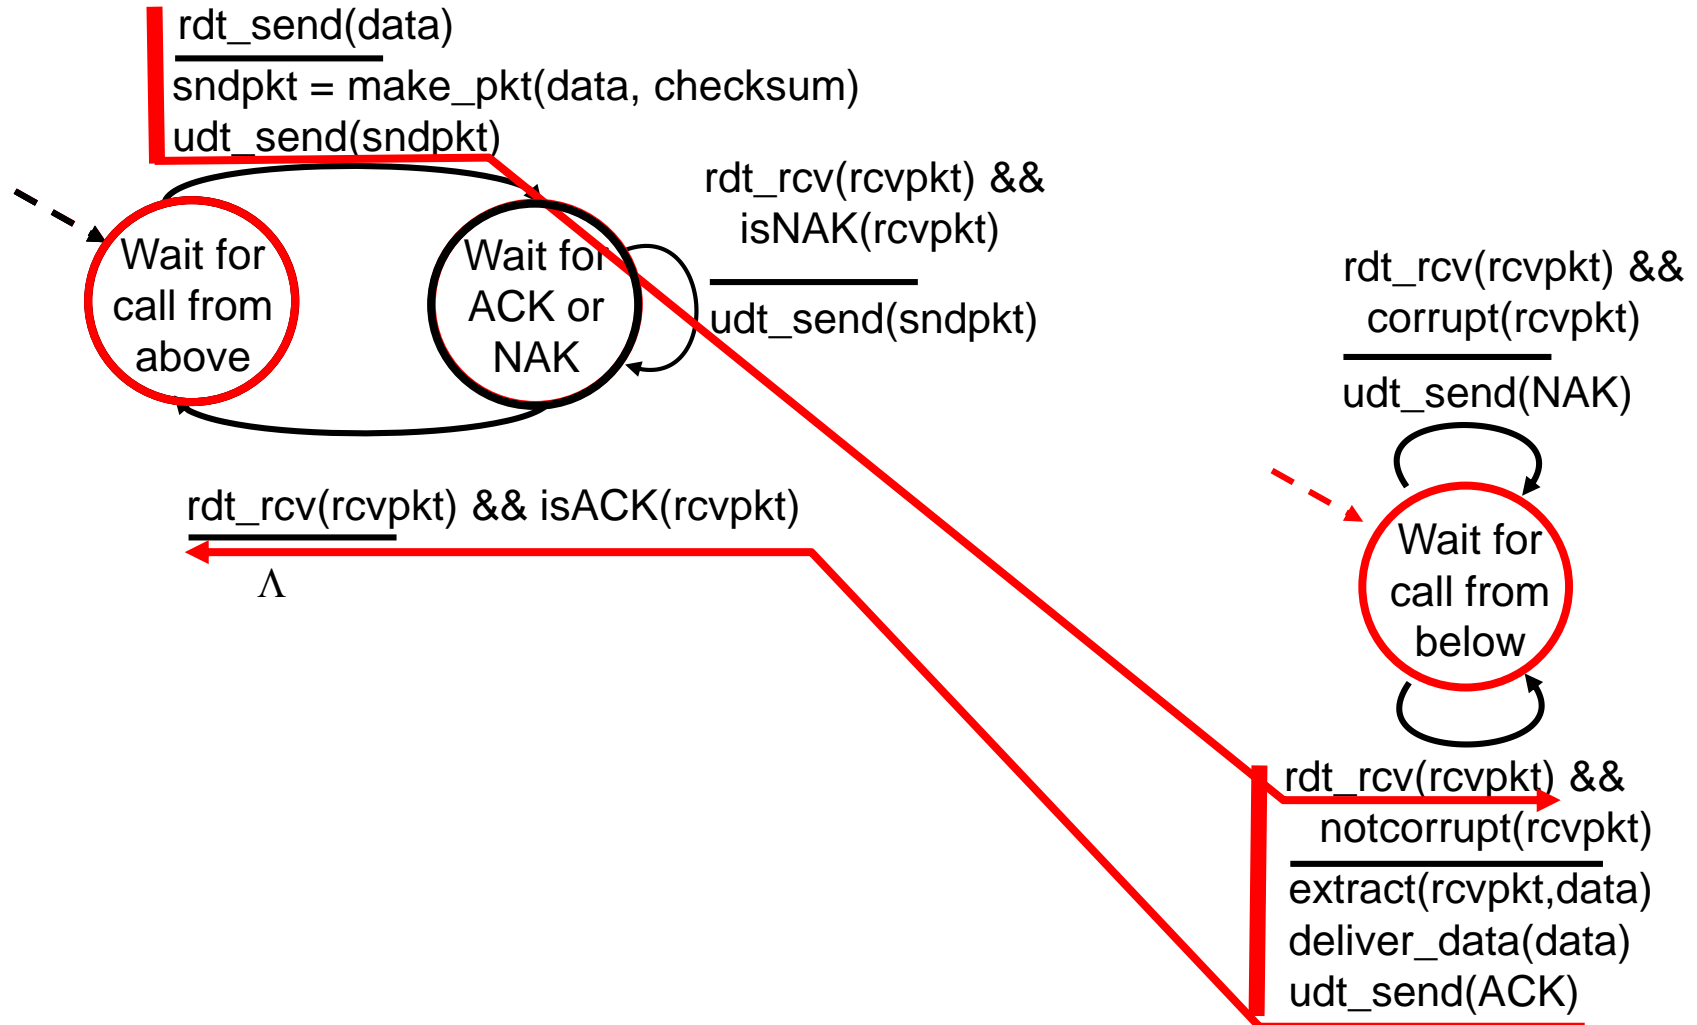


sender

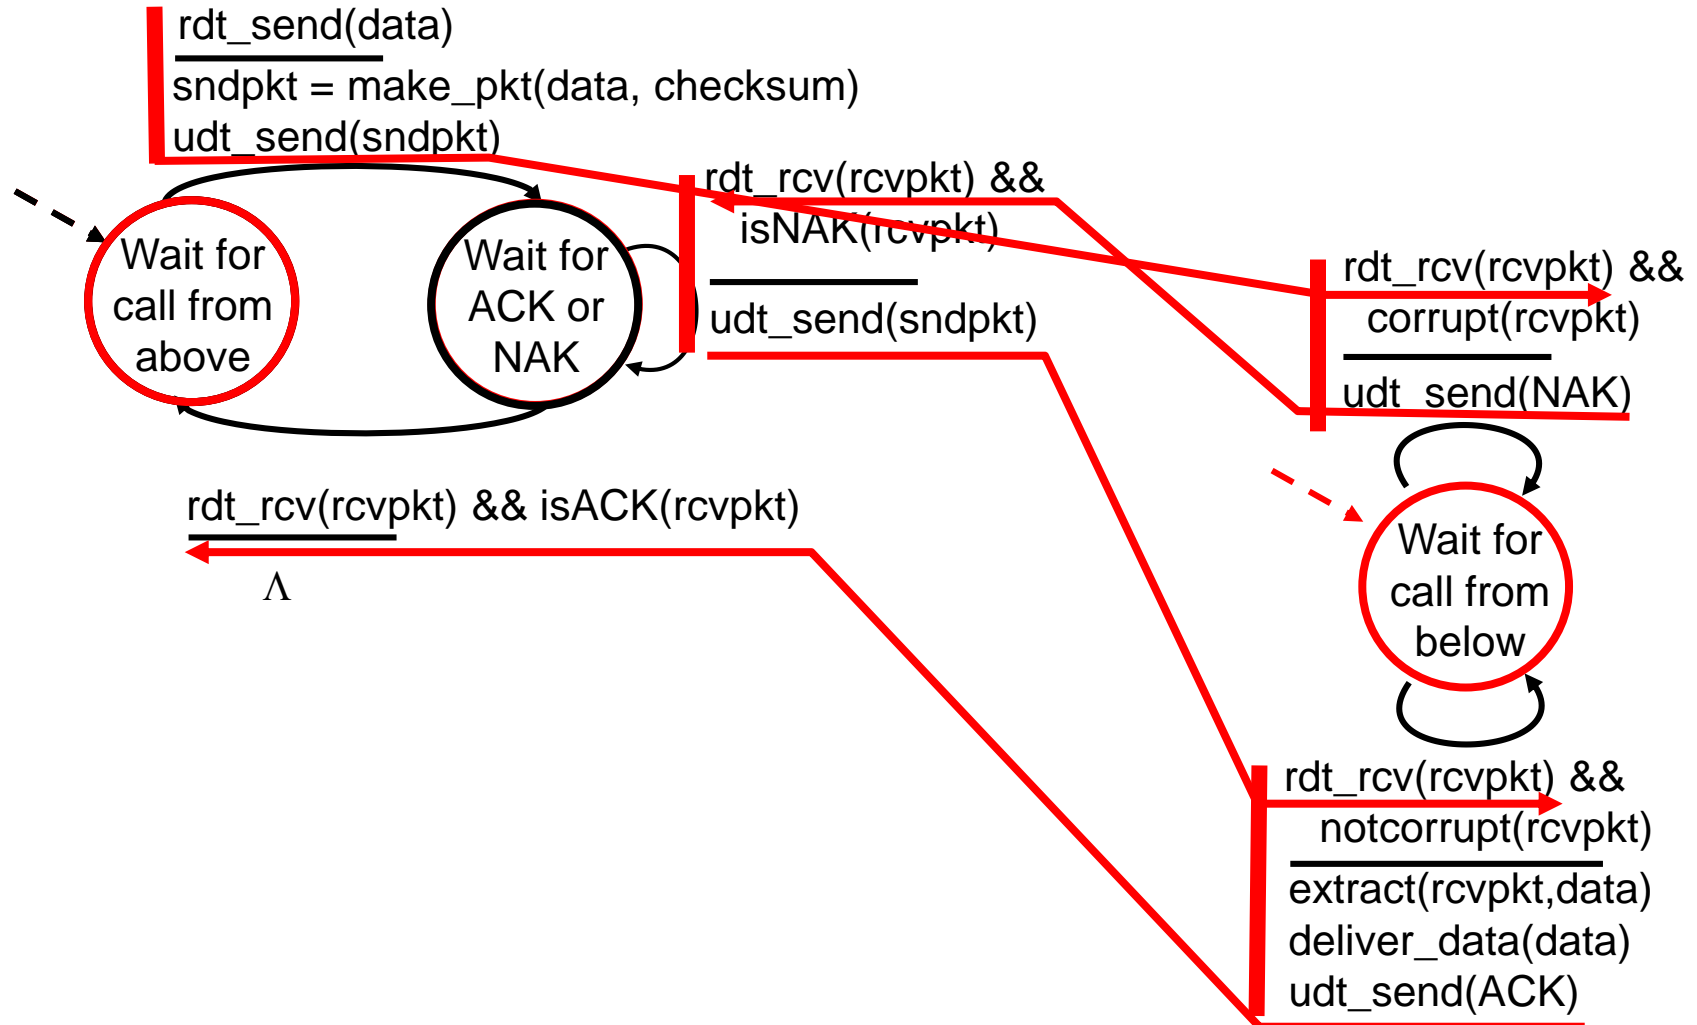
receiver



# rdt2.0: operation with no errors



# rdt2.0: error scenario





# rdt2.0 has a fatal flaw!

## what happens if ACK/NAK corrupted?

- sender doesn't know what happened at receiver!
- can't just retransmit:  
possible duplicate

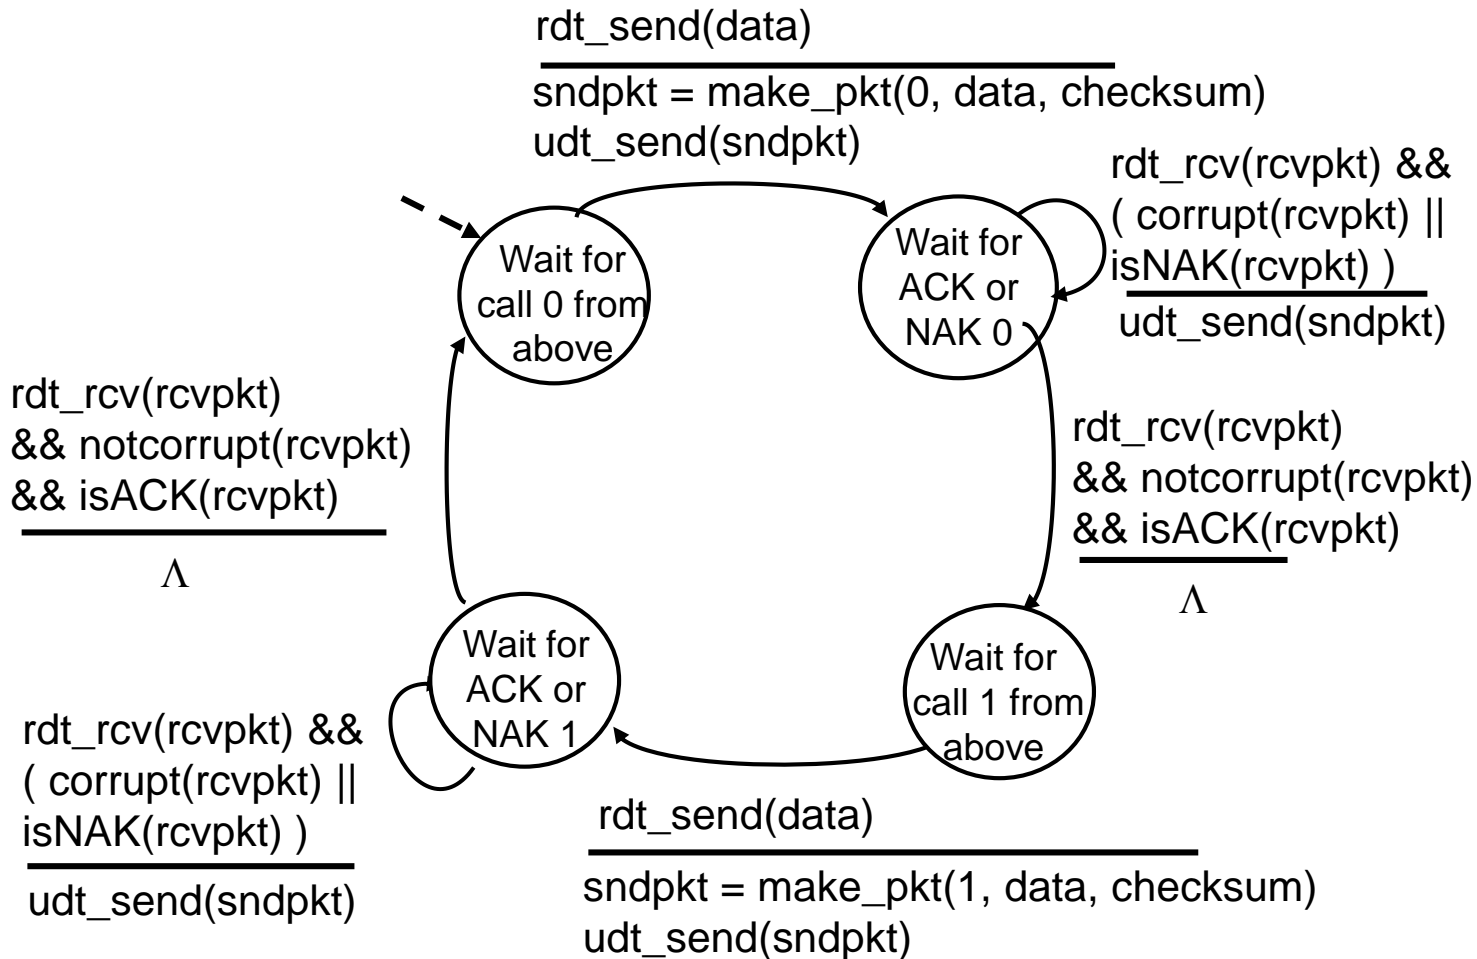
## handling duplicates:

- sender retransmits current pkt if ACK/NAK corrupted
- sender adds *sequence number* (seq) to each pkt
- receiver discards (doesn't deliver up) duplicate pkt

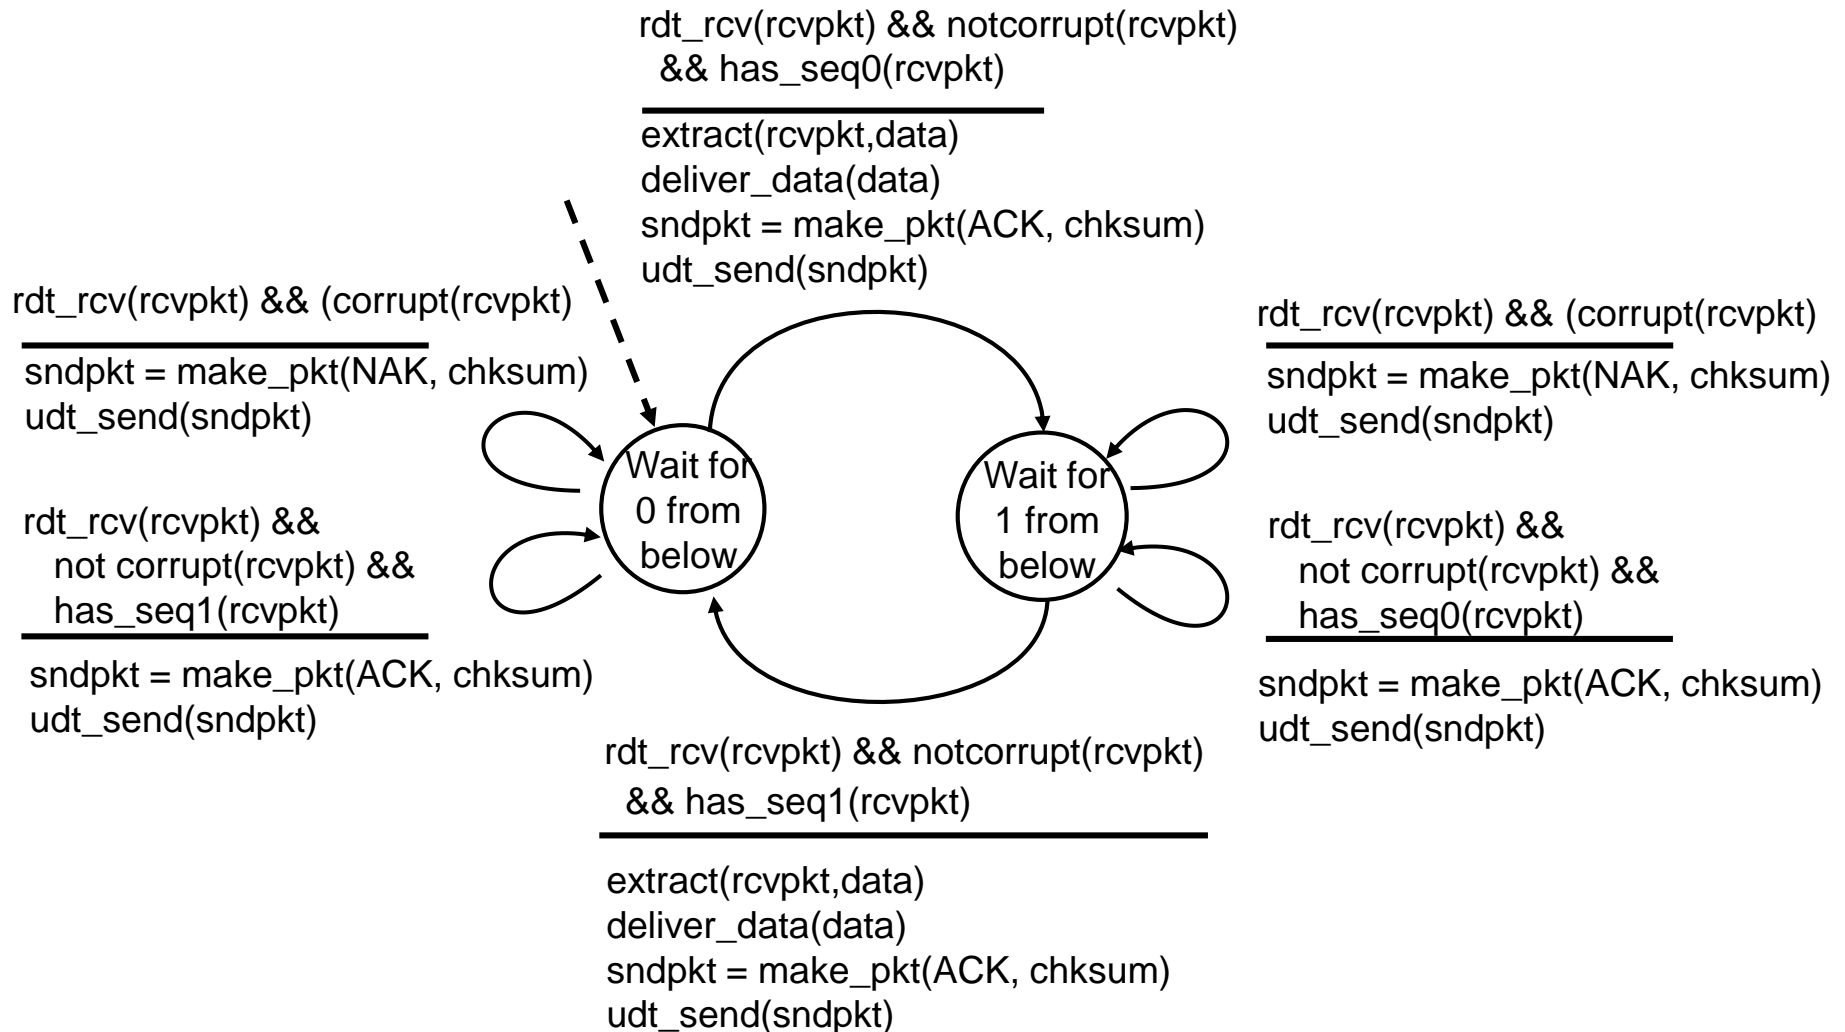
## stop and wait

sender sends one packet,  
then waits for receiver  
response

# rdt2.1: sender, handles garbled ACK/NAKs



# rdt2.1: receiver, handles garbled ACK/NAKs



# rdt2.1: discussion

## sender:

- seq # added to pkt
- two seq. #'s (0,1) will suffice. Why?
- must check if received ACK/NAK corrupted
- twice as many states
  - state must “remember” whether “expected” pkt should have seq # of 0 or 1

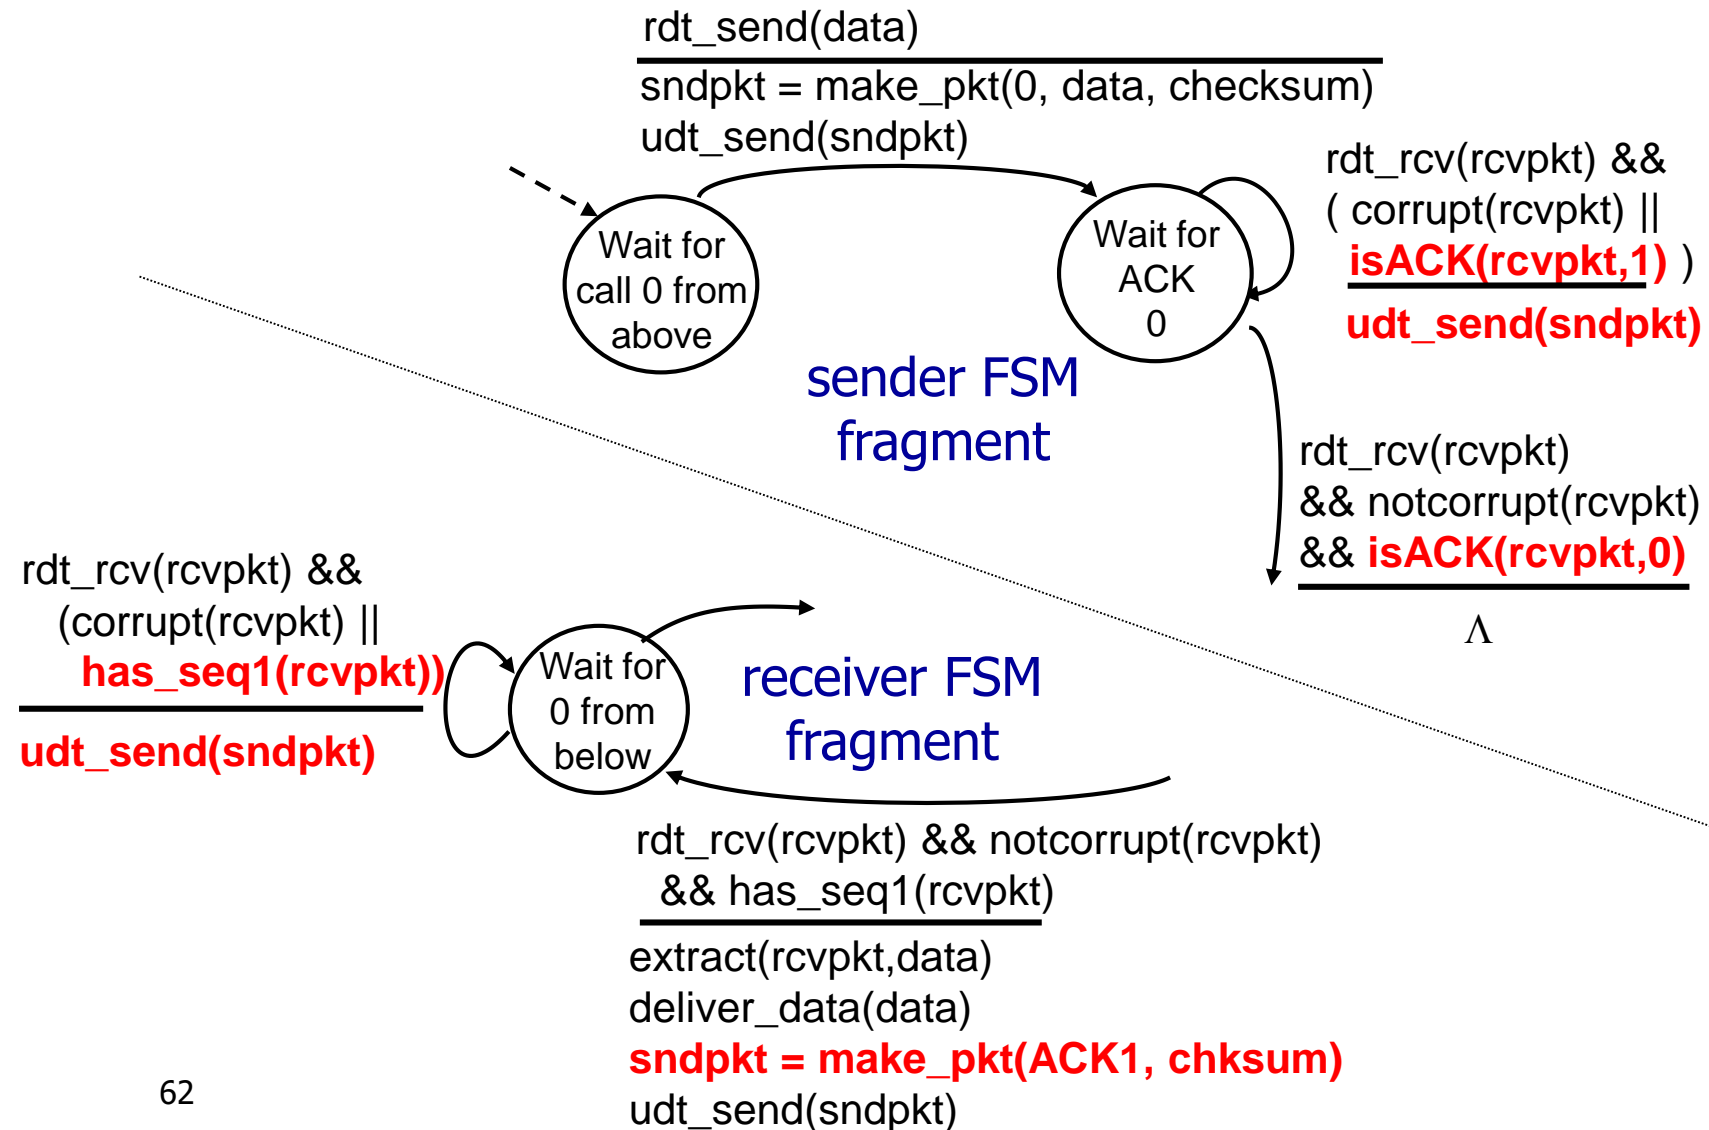
## receiver:

- must check if received packet is duplicate
  - state indicates whether 0 or 1 is expected pkt seq #
- note: receiver can *not* know if its last ACK/NAK received OK at sender

## rdt2.2: a NAK-free protocol

- same functionality as rdt2.1, using ACKs only
- instead of NAK, receiver sends ACK for last pkt received OK
  - receiver must *explicitly* include seq # of pkt being ACKed
- duplicate ACK at sender results in same action as NAK: *retransmit current pkt*

# rdt2.2: sender, receiver fragments



# rdt3.0: channels with errors *and* loss

## new assumption:

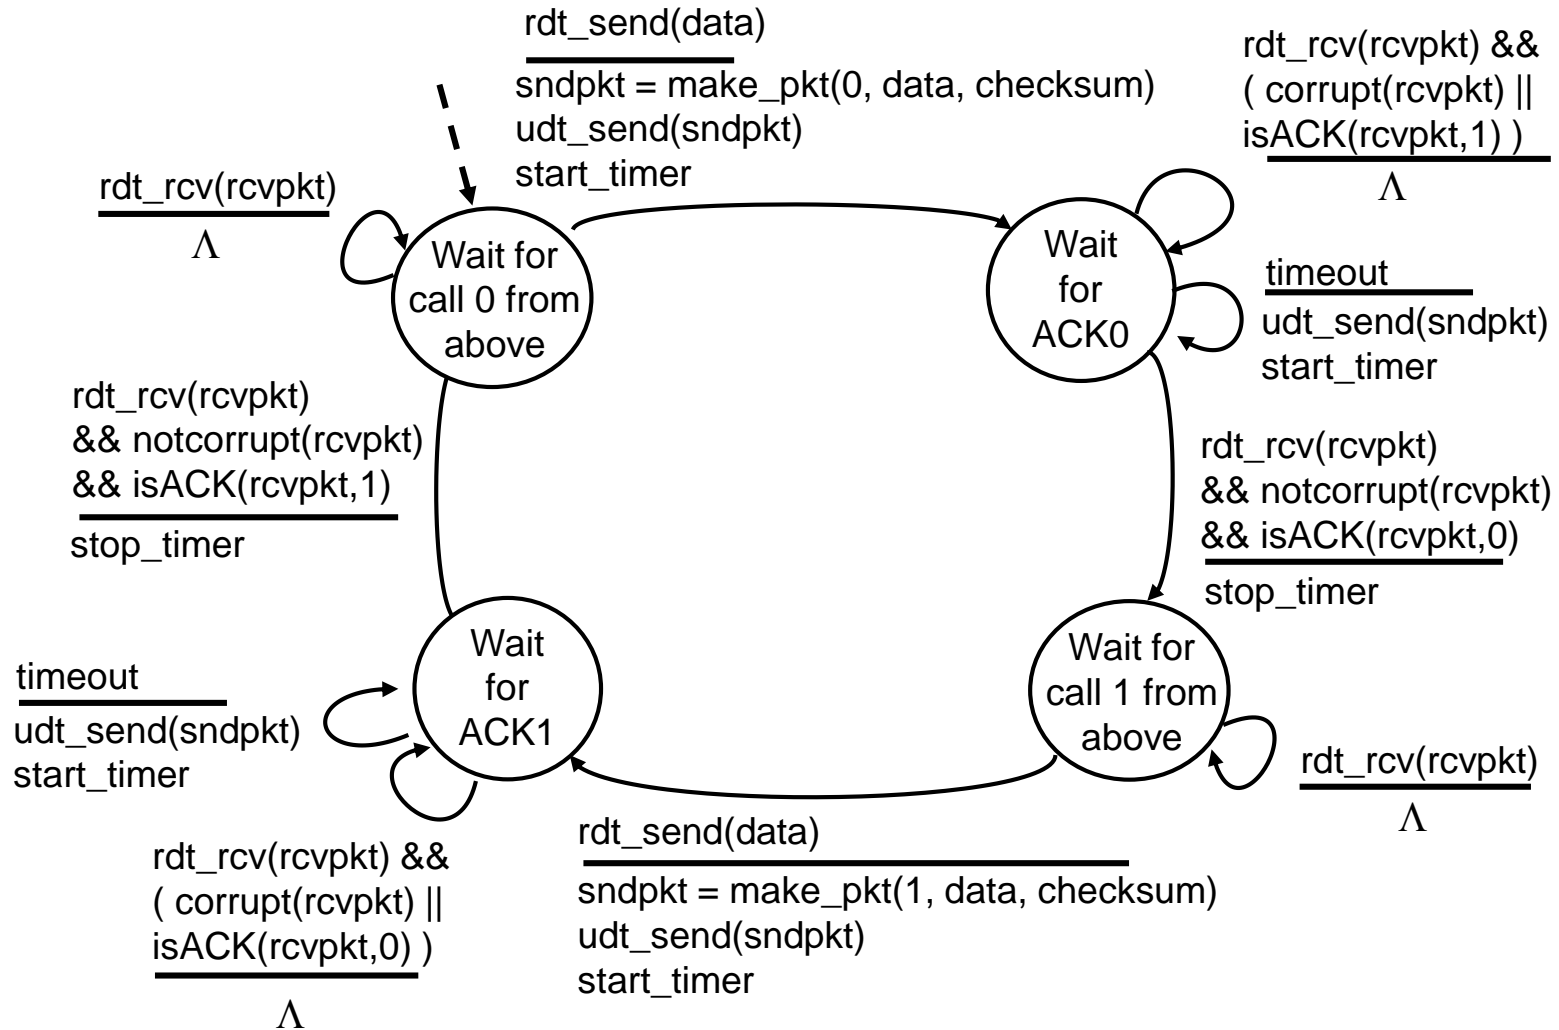
underlying channel can also lose packets (data, ACKs)

- checksum, seq. #, ACKs, retransmissions will be of help ... but not enough

approach: sender waits “reasonable” amount of time for ACK

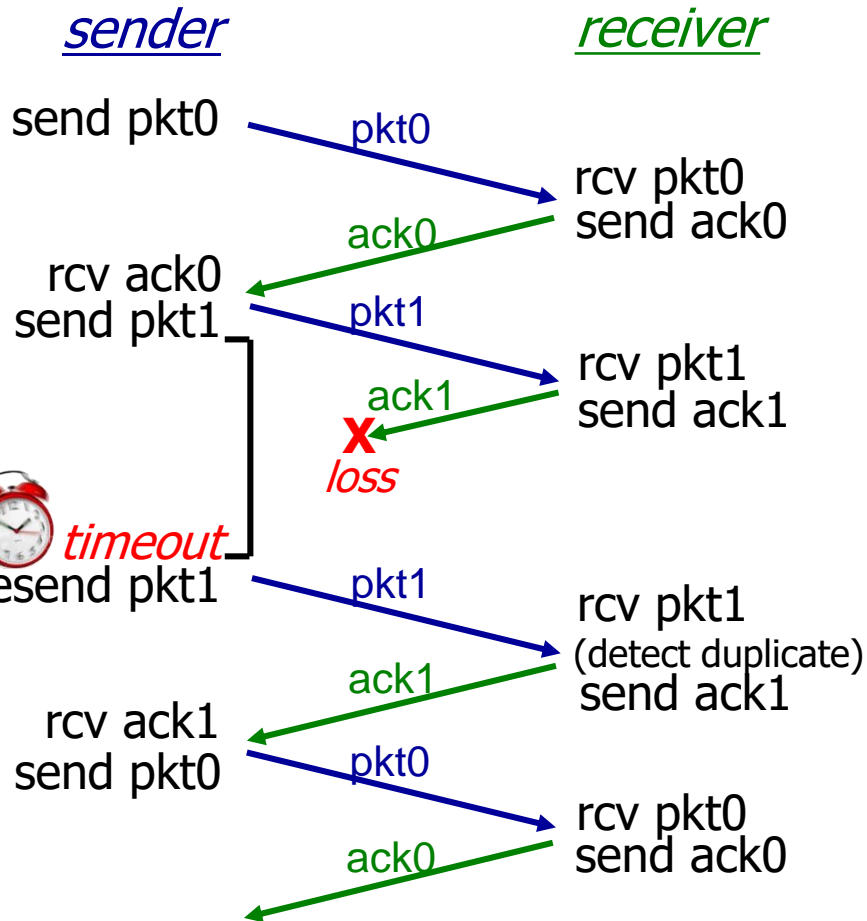
- retransmits if no ACK received in this time
- if pkt (or ACK) just delayed (not lost):
  - retransmission will be duplicate, but seq. #'s already handles this
  - receiver must specify seq # of pkt being ACKed
- requires countdown timer

# rdt3.0 sender

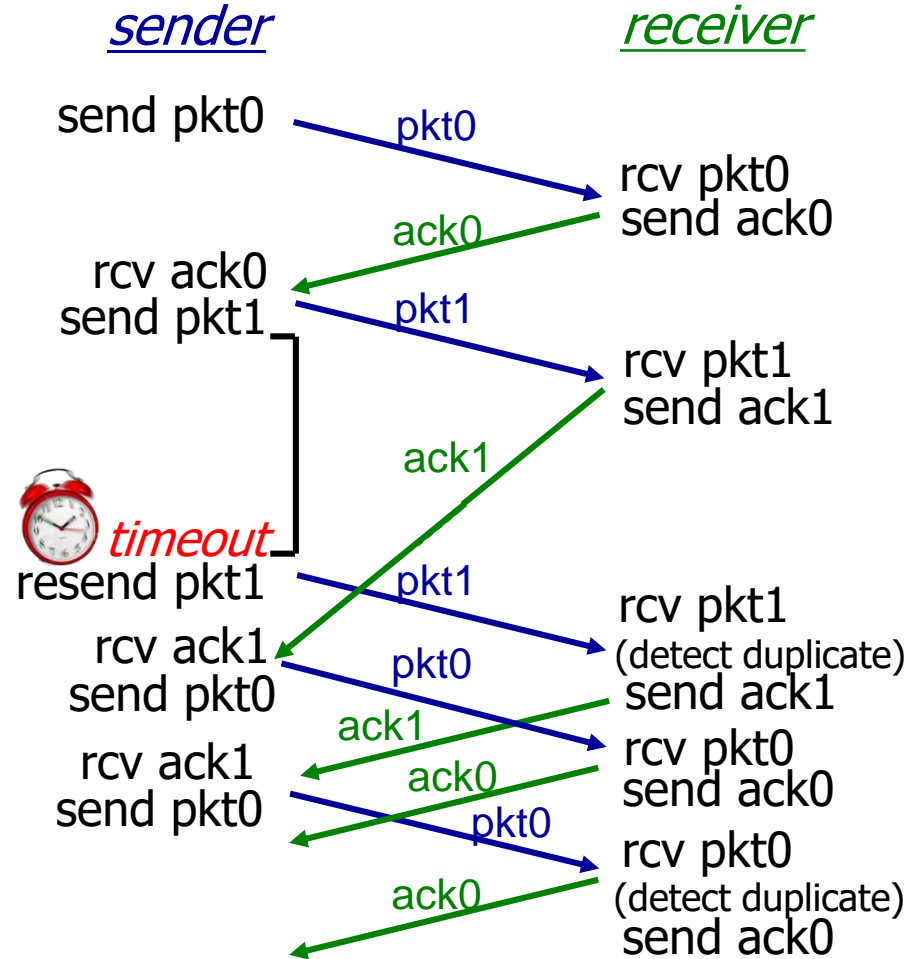




# rdt3.0 in action



(c) ACK loss



(d) premature timeout/ delayed ACK

# What have we learned?

- We are starting on our journey to create “reliable data transfer” over an unreliable network.
- Key concepts:
  - Sequence number – measures which packet we are currently sending
  - ACK (acknowledgement) – numbered, says that a specific packet has definitely been arrived
  - Timeout – a time to wait before we think a packet is lost

# Transport layer summary

- So far:
  - 3.1 transport-layer services
  - 3.2 multiplexing and demultiplexing
  - 3.3 connectionless transport: UDP
  - 3.4 principles of reliable data transfer
- Next week of lectures:
  - 3.4 principles of reliable data transfer (continued)
  - 3.5 connection-oriented transport: TCP
  - 3.6 principles of congestion control
  - 3.7 TCP congestion control